

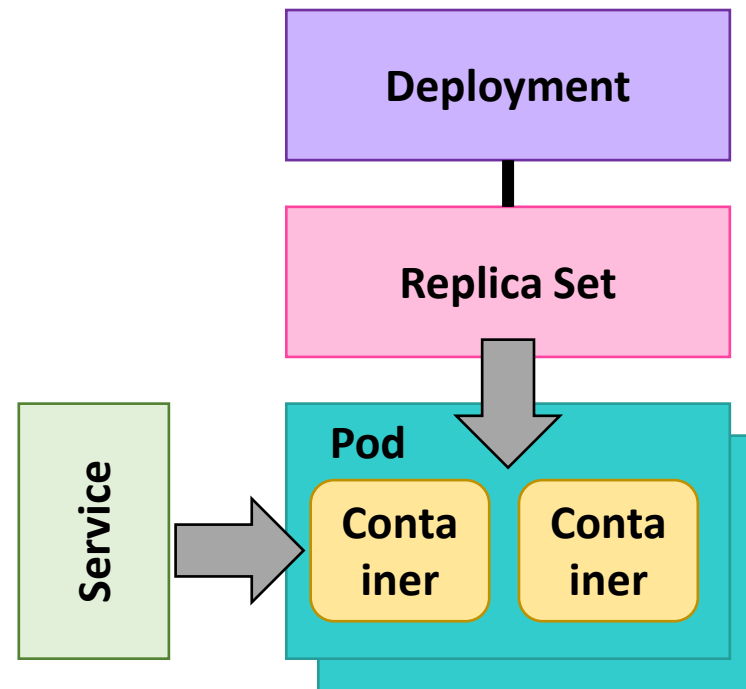


Kubernetes

Part 2



Kubernetes





Volumes

- Volumes are storage that are shared by containers in a Pod
 - Allocated by the Pod, usually a shared directory in the Pod
 - Not visible outside of Pod
- Tied to the lifecycle of a Pod viz. its removed when the Pod is delete
 - Unlike Docker volumes where they are durable
- Different types of volumes
 - Eg. hostPath, NFS, iSCSI, fibre channel, empty directory, etc.
- **hostPath** and **emptyDir** type is good for sharing data between containers in a Pod
 - Eg. The example of file puller and web server


emptyDir will be deleted when pod is deleted



hostPath is similar to the concept of docker bind mount

Defining a Volume

Same syntax for creating for Pod templates in deployment objects





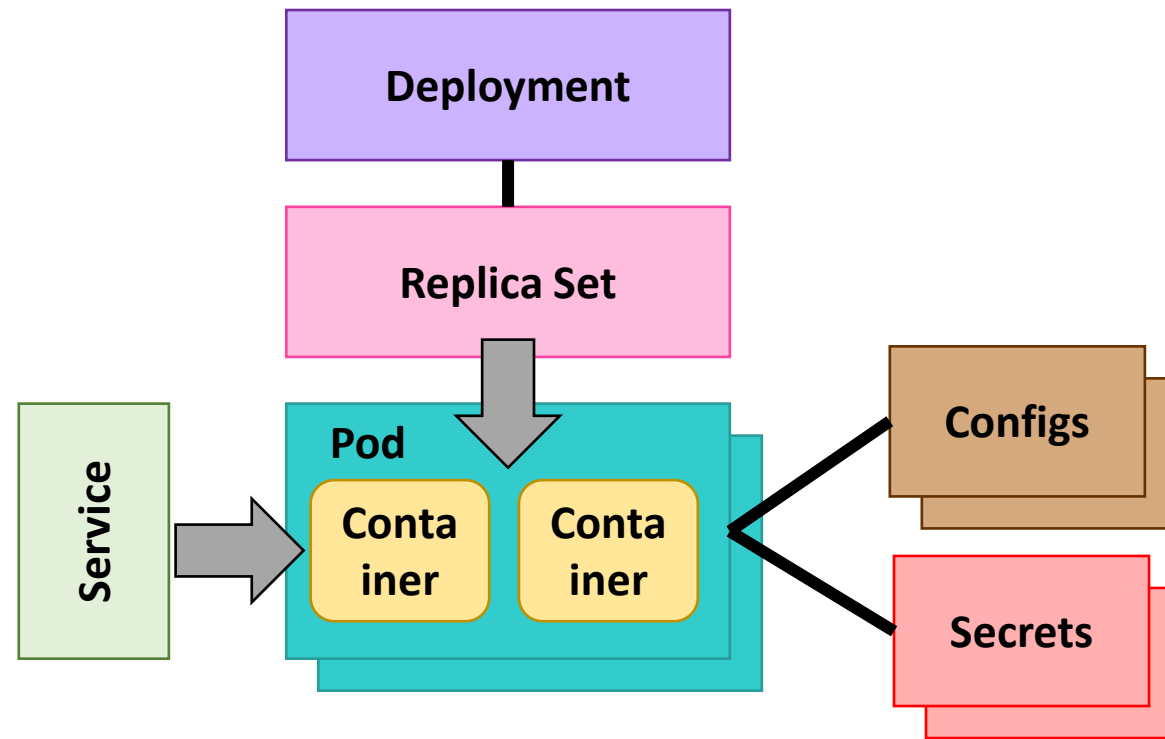
Mounting ConfigMaps and Secrets

- `configMap` and `secrets` can be mounted as volumes
 - Keys becomes the filename, the value is the file content
- Use cases
 - Passing configuration files
eg `nginx.conf`

```
spec:
  volumes:
    - name: html-vol
      configMap:
        name: html-assets
  containers:
    - name: nginx
      image: nginx
      volumeMounts:
        - name: html-vol
          mountPath: /usr/share/nginx/html
          readOnly: true
```



Kubernetes





Persistent Storage

- Kubernetes can dynamically provision storage
 - Eg. User ask for 50GB volume to caching images
- Kubernetes allows storage to be either statically or dynamically provisioned
 - Static provision - an administrator will need to first provision the storage manually
 - Dynamic provision - the user describes the type of storage that is required; Kubernetes will attempt to provision based on the user's requirements
- Once a persistent storage has been allocated and claimed/reserved, a Pod can mount the volume like any regular volume
- Persistent volumes lifecycle are not tied to the Pod's lifecycle
 - Unlike volumes, persistent volumes will not be deleted when a Pod is deleted
 - This behaviour can be configured

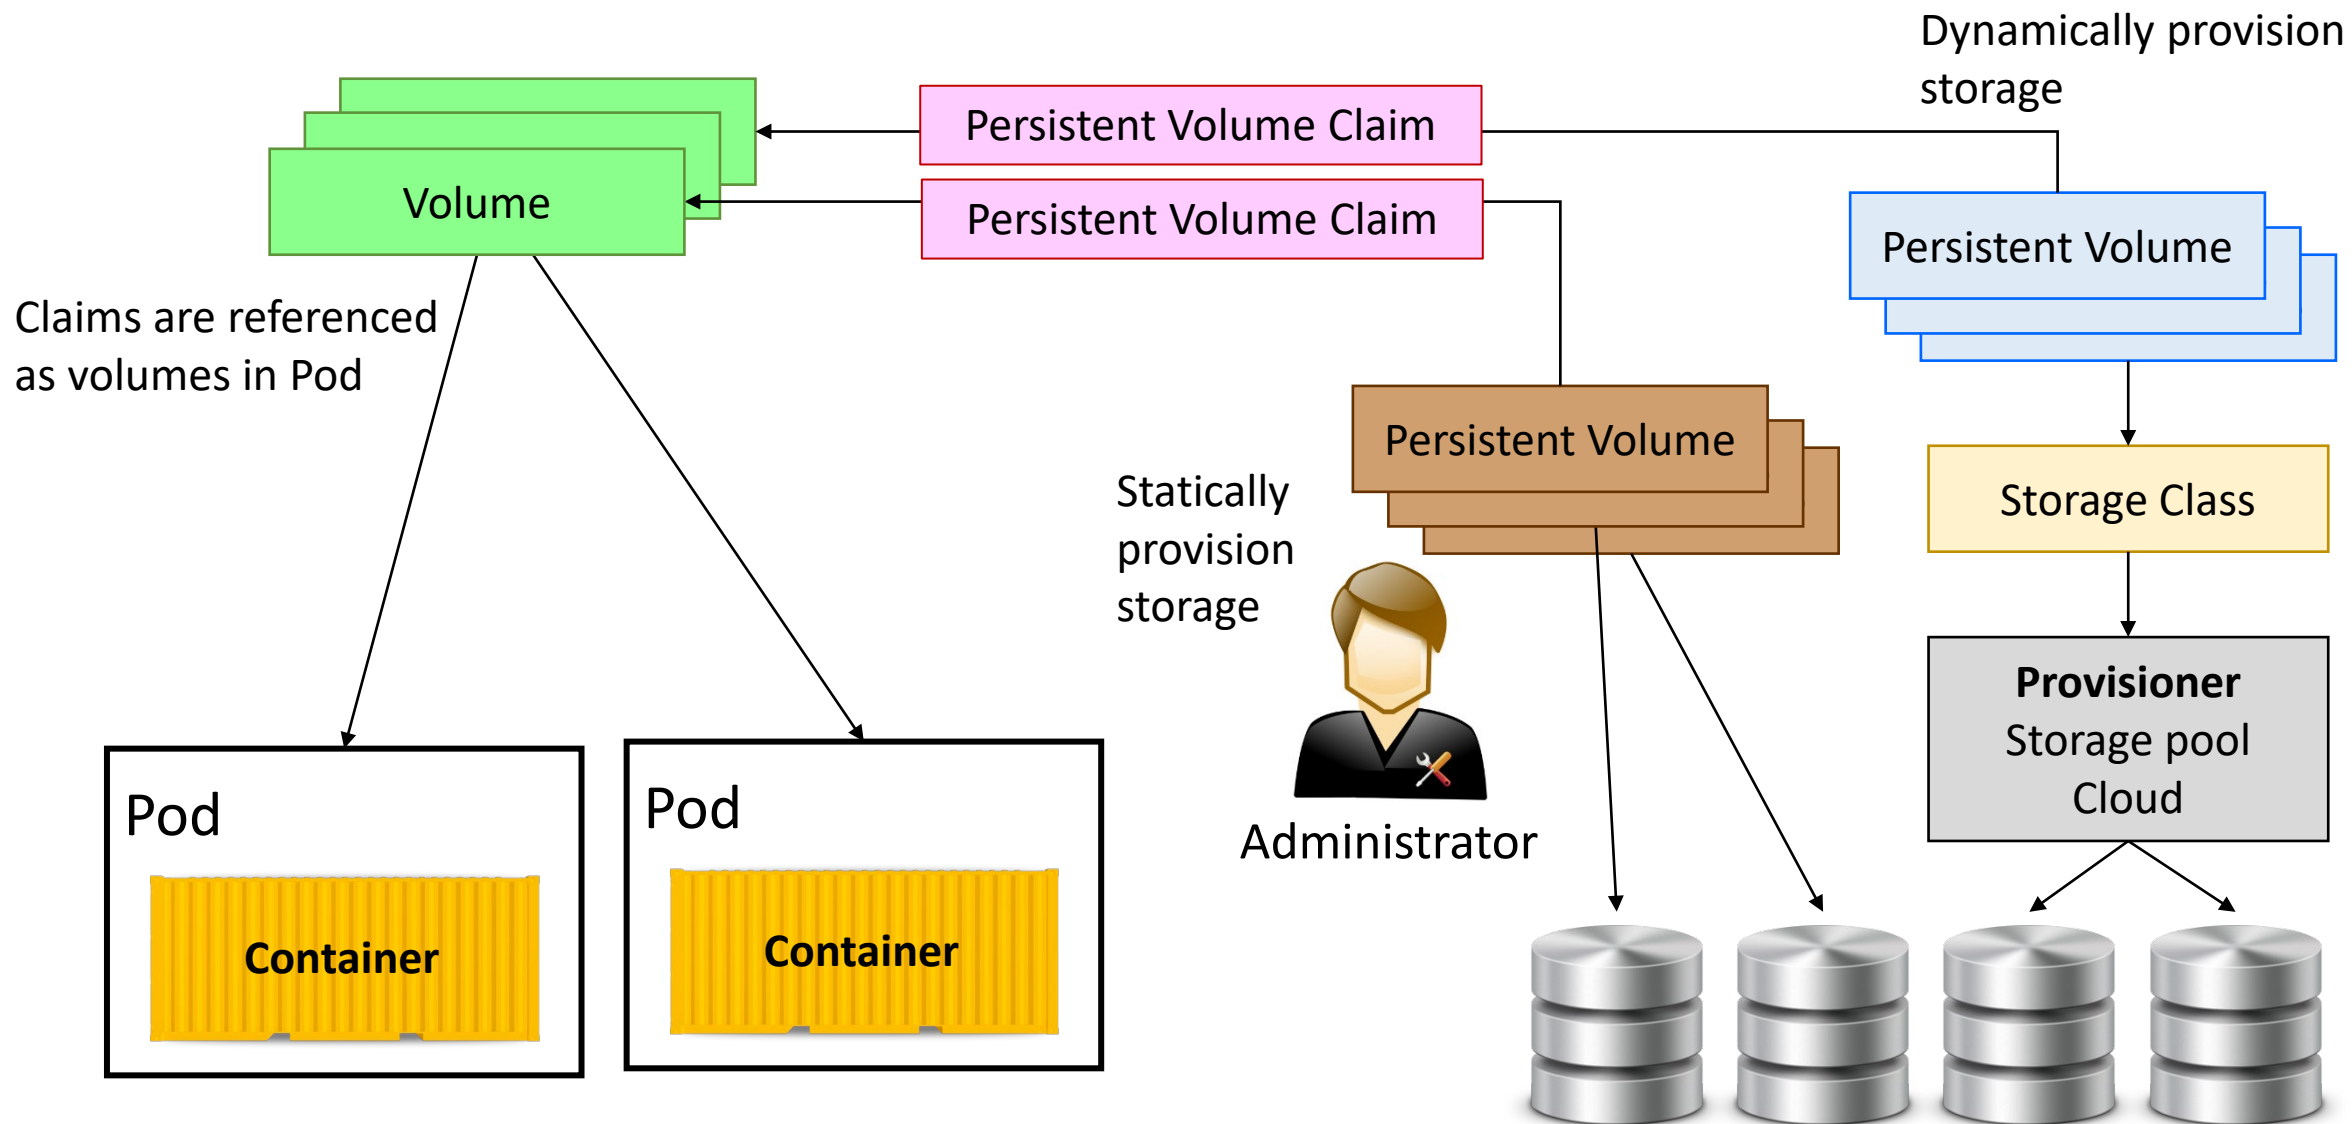


Key Concepts

- Storage class - a type of storage
 - Who the provisioner, storage specific details, retention policy, etc.
- Persistent volume - the actual storage
 - A piece of storage provisioned by an administrator or thru storage class
 - Supports may different storage type
 - AWS EBS, Azure File Service, Cinder, fibre channel, GCP Disk, NFS, etc.
 - Different type of access mode - exclusive or shared
- Persistent Volume claim - when a persistent volume has been allocated for use, the volume is staid to be claimed



Persistent Volume





Static vs Dynamic

Static

- Administrator has to manually allocate storage and map it to a persistent volume
- Users can then claim this volume

Dynamic

- When Kubernetes tries to resolve a claim and the persistent volume is unavailable
- It looks for a storage class that best matches the request storage
- Dynamically creates the persistent volume using the provisioner



Defining a Persistent Volume Claim

```
apiVersion: v1
```

```
kind: PersistentVolumeClaim
```

```
meta-data:
```

```
  name: myapp-pvc
```

```
spec:
```

```
  accessModes:
```

```
    - ReadWriteOnce
```

List of access modes

ReadOnlyMany

ReadWriteMany

```
  resources:
```

```
    requests:
```

```
      storage: 5Gi
```

```
  storageClassName: standard
```

Get storage class name(s) with
kubectl get storageclass



Mounting a Persistent Volume

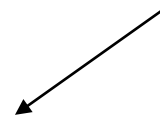
```
apiVersion: v1
kind: Pod
meta-data:
  name: myapp
```

```
spec:
```

```
  volumes:
```

```
    - name: data-volume
      persistentVolumeClaim:
        claimName: myapp-pvc
```

Specify the claim name



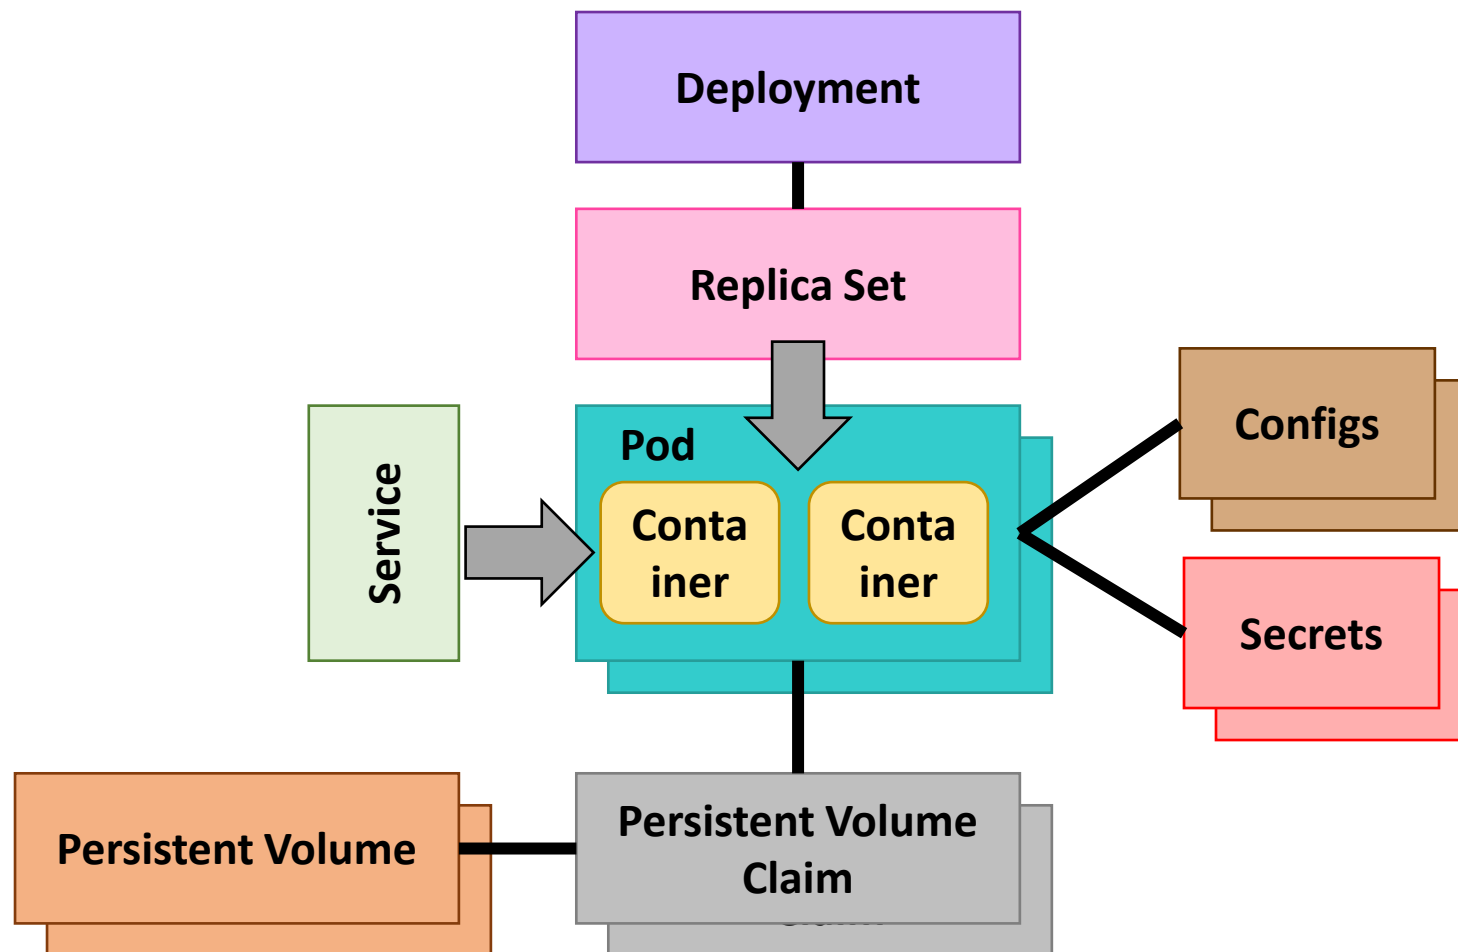
```
  containers:
```

```
    - name: myapp
      volumeMounts:
        - mountPath: /app/public
          name: data-volume
```

```
    ...
```



Kubernetes





Persistence Volume Management

- Display persistence volume detail
 - Persistence volume - `kubectl get pv`
 - Persistence volume claim - `kubectl get pvc`
 - Storage classes - `kubectl get sc`
- Delete persistence volume


```
kubectl delete pvc <name>
```

```
kubectl delete pv <name>
```



Load Balancer and Ingress

- By default services are allocated a cluster IP
 - Only accessible within the cluster
- Load balancer exposes the service to the public
 - Accessible from outside of the cluster
 - Load balancer will redirect the request to pods based on its routing policy
 - Another way to allow external access is via node port
- Load balancer are resources that are provisioned from the underlying cloud platform
 - May have more features that you require
 - Also cost more





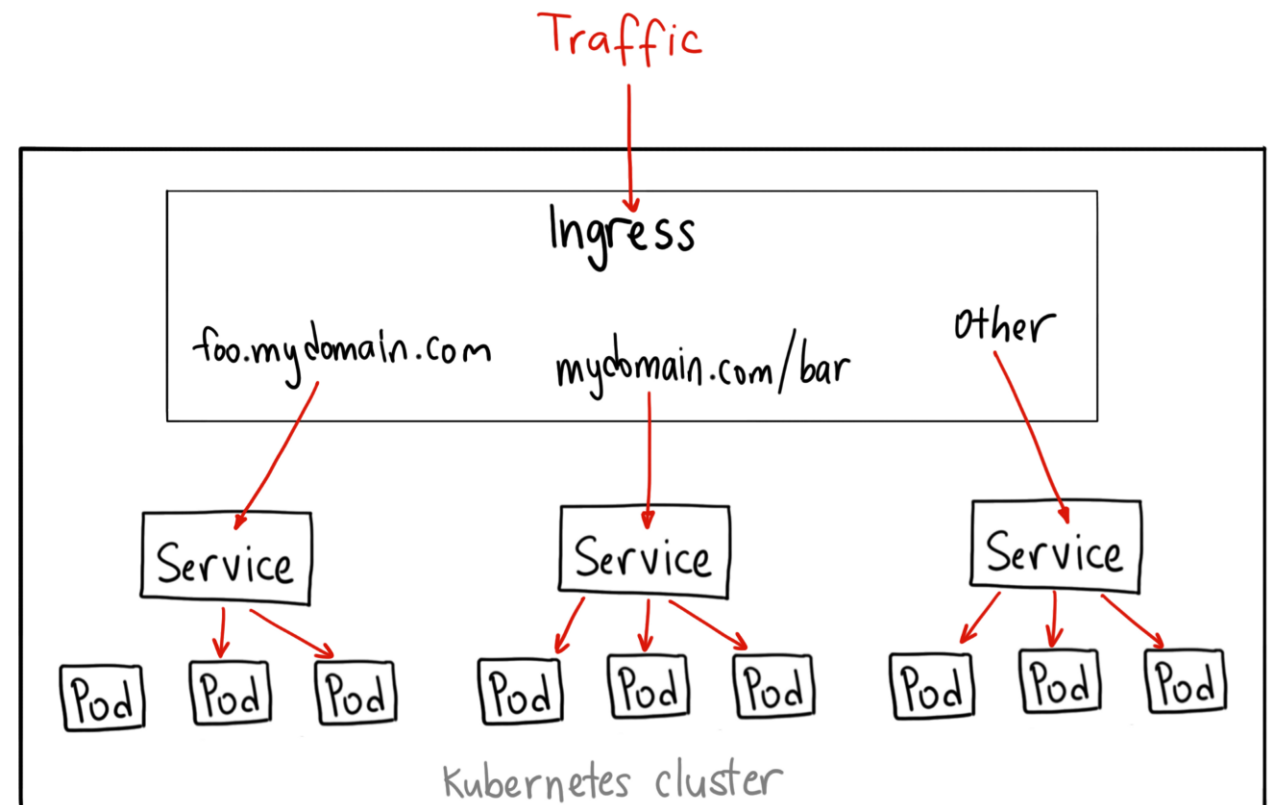
Load Balancer and Ingress

- An ingress, like a LoadBalancer service, allows traffic into the cluster
 - Provisions an 'external' load balancer
 - Shared by many services within the cluster
 - More cost effective
- May require to provision an Ingress controller
 - May not be available
- NGINX Ingress controller is a popular ingress controller
 - Deploys NGINX as Ingress
 - <https://github.com/kubernetes/ingress-nginx>



Ingress

- Application layer (L7) router that sits in front of multiple services
- Define a set of routing rules on how services are access externally
 - Eg. 2 services, one for search one for checkout. Might map to /search and /checkout
- Rules are applied to ingress controllers which performs the actual routing
 - Controllers might be a cloud provider's load balancer or Nginx reverse-proxy





Defining an Ingress

```
apiVersion: networking.k8s.io/v1
```

```
kind: Ingress
```

```
metadata:
```

```
  name: myapp
```

```
  annotations:
```

```
    nginx.ingress.kubernetes.io/rewrite-target: "/"
```

```
    nginx.ingress.kubernetes.io/ssl-redirect: "false"
```

```
spec:
```

```
  ingressClassName: nginx
```

```
  rules:
```

```
    - host: acme.com
```

```
      http:
```

```
        paths:
```

```
          - path: /hello
```

```
            pathType: Prefix
```

```
            backend:
```

```
              service:
```

```
                name: myapp
```

```
                port:
```

```
                  number: 8080
```

Change/rewrite a matched resource
to its root e.g /hello to /

Used to configure
NGINX ingress
controller

One or more of
these rules to
specify which
services to handle
what resource

Select the ingress controller to use

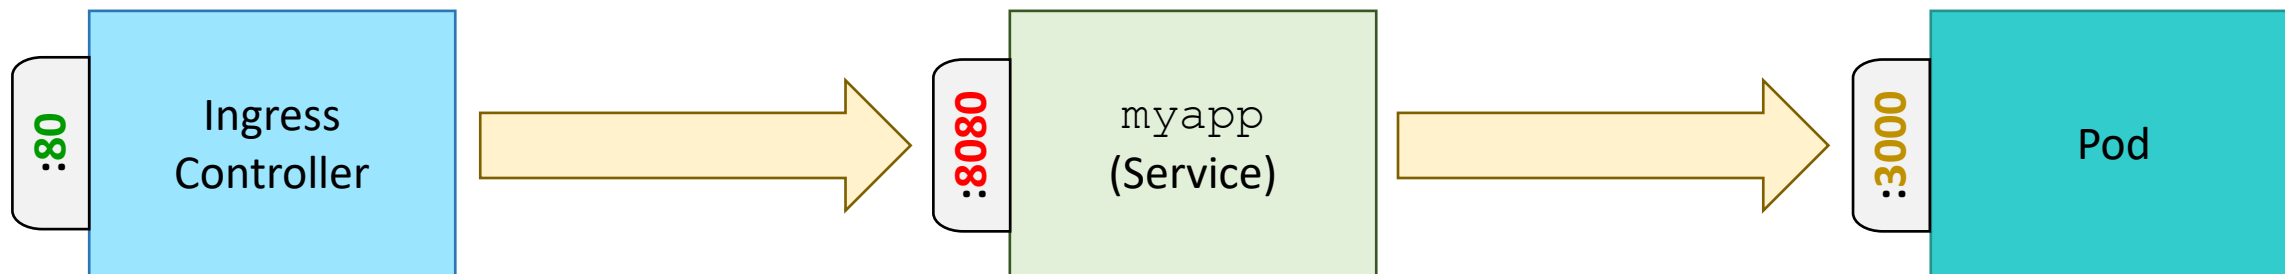


Ingress Ports

```
kind: Ingress
spec:
  rules:
  - host: acme.com
    http:
      paths:
      - path: /
        pathType: Prefix
        backend:
          service:
            name: mysvc
            port:
              number: 8088
```

```
kind: Service
spec:
  ports:
  - port: 8080
    targetPort: 3000
```

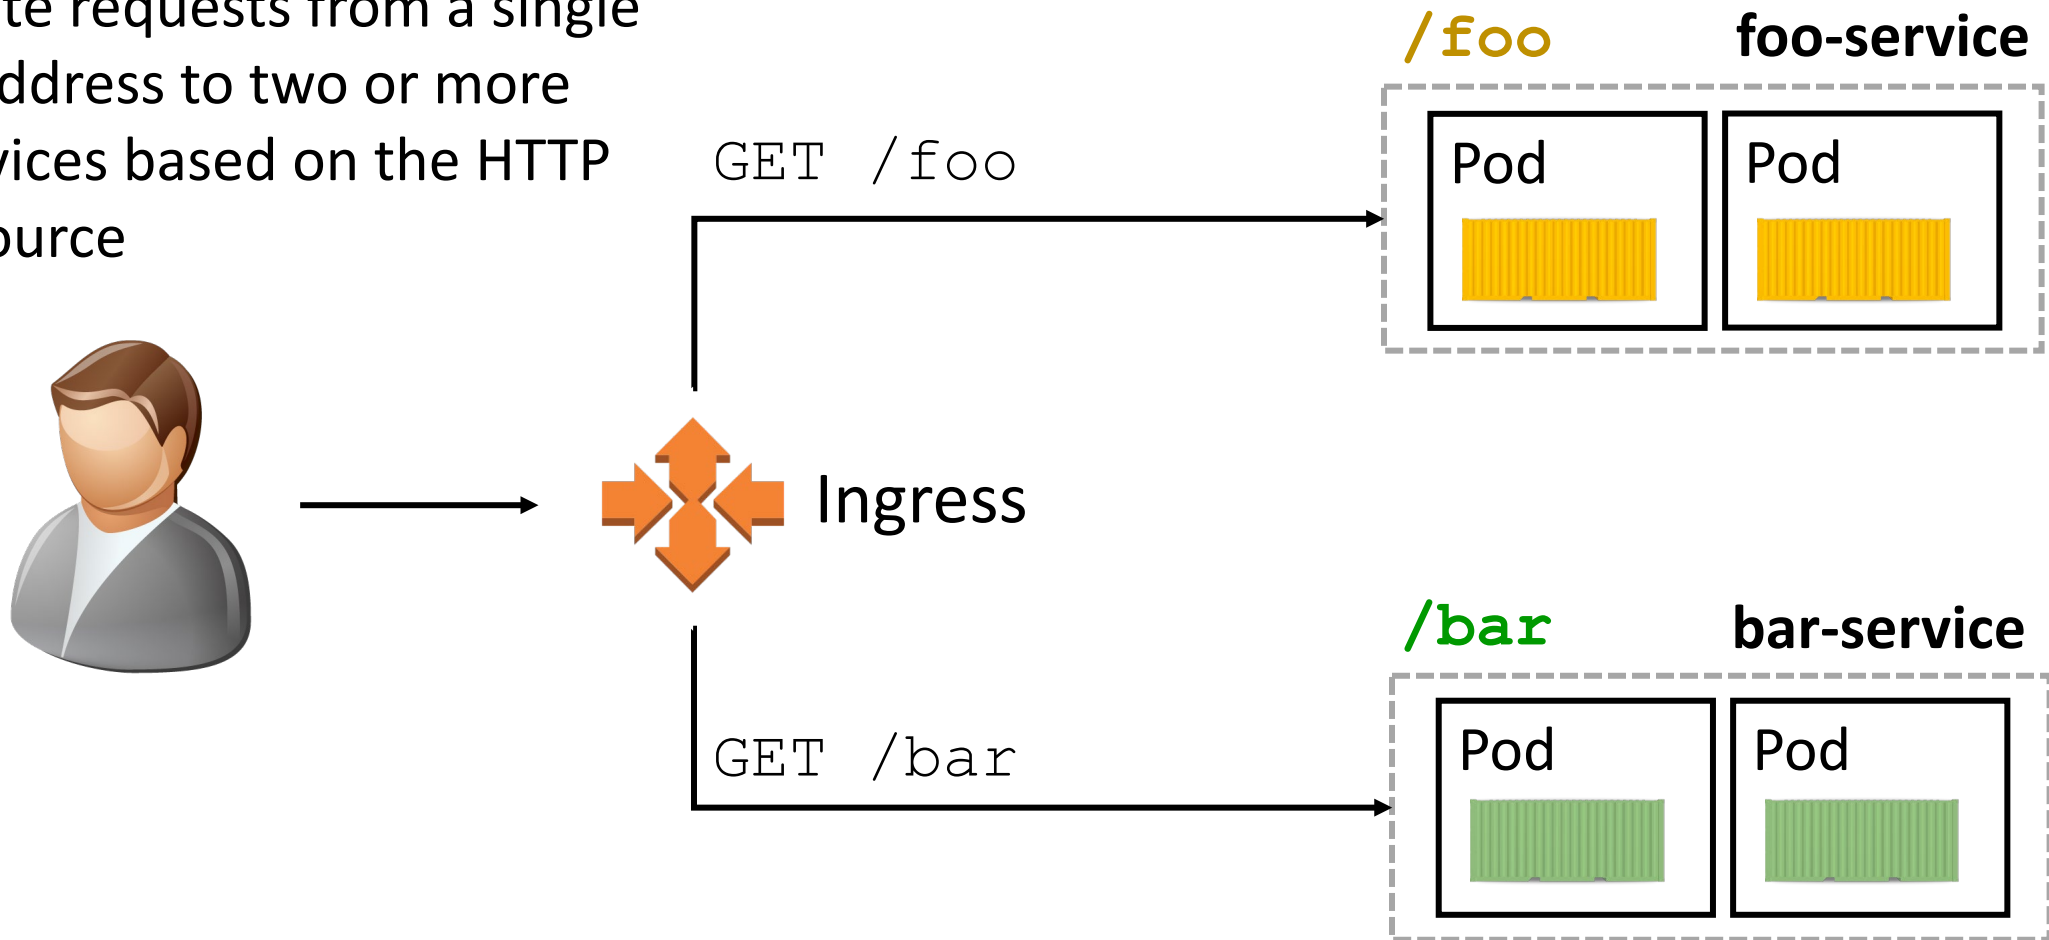
```
kind: Deployment
spec:
  containers:
  - ports:
    - containerPort: 3000
```





Ingress - Fan Out

Route requests from a single IP address to two or more services based on the HTTP resource





Ingress Fan Out Example

```
apiVersion: networking.k8s.io/v1
kind: Ingress
metadata:
  name: myapp-ingress
  annotations:
    nginx.ingress.kubernetes.io/rewrite-target: /
spec:
  ingressClassName: nginx
  rules:
  - host: acme.com
    http:
      paths:
      - path: /foo
        pathType: Prefix
        backend:
          service:
            name: foo-service
            port:
              number: 8000
      - path: /bar
        pathType: Prefix
        backend:
          service:
            name: bar-service
            port:
              number: 8001
```

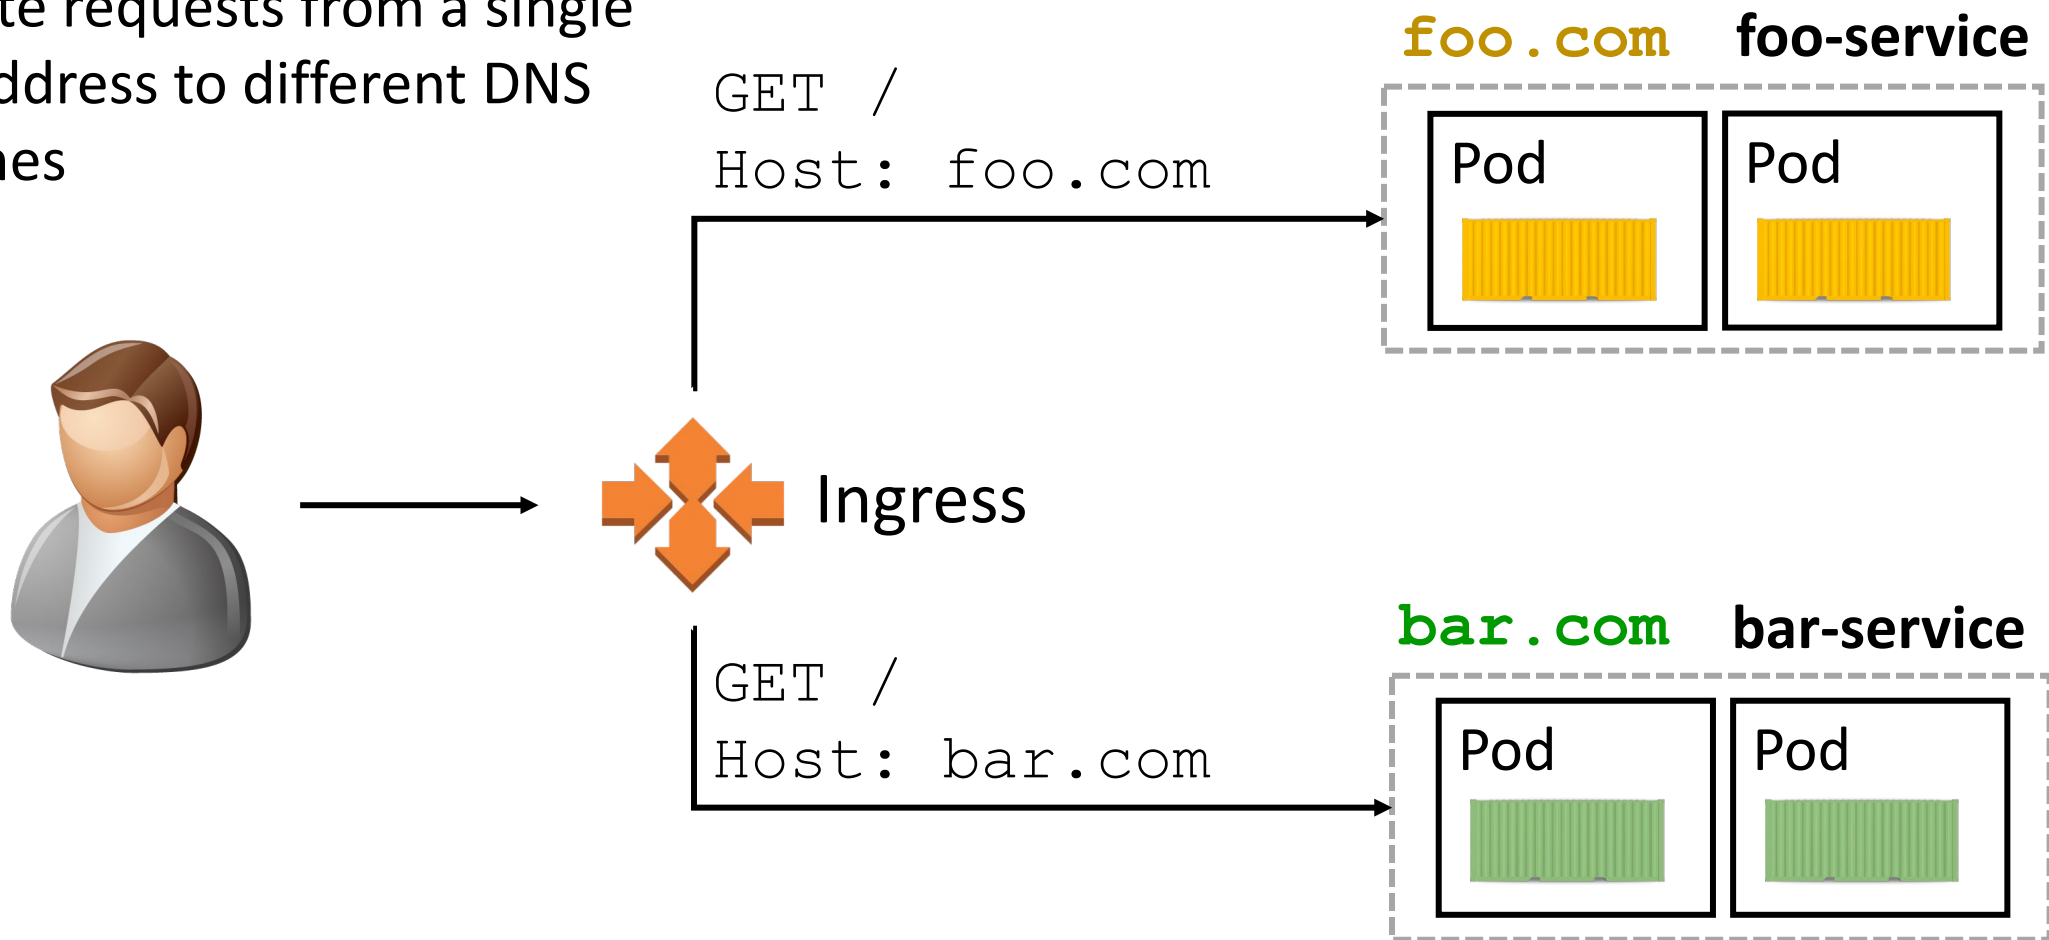
Note: If the 2 services are web application, then all resources references (eg in HTML) are now rooted under `/foo` or `/bar`
So the web application must take this into account
One option is to use relative reference or use `<base>`

Request is routed to these 2 services
`acme.com/foo` or `acme.com/bar`



Ingress - Virtual Host

Route requests from a single IP address to different DNS names





Ingress Virtual Host Example

```
apiVersion: networking.k8s.io/v1
```

```
kind: Ingress
```

```
metadata:
```

```
  name: myapp-ingress
```

```
spec:
```

```
  ingressClassName: nginx
```

```
  rules:
```

```
    - host: foo.com
```

```
      http:
```

```
        paths:
```

```
          - pathType: Prefix
```

```
            backend:
```

```
              name: foo-service
```

```
              port:
```

```
                number: 8080
```

```
    - host: bar.com
```

```
      http:
```

```
        paths:
```

```
          - pathType: Prefix
```

```
            backend:
```

```
              name: bar-service
```

```
              port:
```

```
                number: 8080
```

foo.com

host

bar.com

host

Request is routed to these 2 services depending on the `Host` attribute.



Kubernetes Annotations

- Annotations are additional/proprietary metadata passed to a controller
 - Usually configuration information
- Allow additional configuration not part of the spec:
 - Eg. ingress resources does not support canary deployment. Can configure canary if using ingress-nginx thru annotations
- <https://kubernetes.github.io/ingress-nginx/user-guide/nginx-configuration/annotations>



ingress-nginx Features

```
apiVersion: networking.k8s.io/v1
```

```
kind: Ingress
```

```
metadata:
```

```
  name: myapp-ingress
```

```
  annotations:
```

```
    nginx.ingress.kubernetes.io/rewrite: /$2
```

```
spec:
```

```
  ingressClassName: nginx
```

```
  rules:
```

```
    - host: acme.com
```

```
      http:
```

```
        paths:
```

```
          - path: /foo(/|$) (.*)
```

```
            pathType: Prefix
```

```
            backend:
```

```
              ...
```

```
          - path: /bar(/|$) (.*)
```

```
            pathType: Prefix
```

```
            backend:
```

Remove the first segment of the resource from the request

Original URL	Rewrote URL
/foo	/
/foo/debug	/debug
/foo/api/customer/123	/api/customer/123



Handling Errors

- If no rules matches the incoming request, then the traffic can be routed to a default backend if it is configured
 - The default backend service can be any of your application
 - Eg. Help page, chatbot

Traffic will be routed to this service if no rule matches
If the default backend is not specified, then will fallback to the ingress controller

```
apiVersion: v1
kind: Ingress
metadata:
  name: myapp-ingress
  annotations:
    nginx.ingress.kubernetes.io/rewrite: /$2
spec:
  defaultBackend:
    service:
      name: defaultBackendService
      port:
        number: 8080
  rules:
  - host: acme.com
    http:
      ...
```

A curved arrow originates from the text "Traffic will be routed to this service if no rule matches" and points directly to the **defaultBackend** field in the YAML code block.



Handling Errors

- An alternative is to install a global default service
 - This feature is specific to the ingress controller
- For `stable/nginx-ingress`, this is done during deployment
 - With helm
- Can only have a single default backend

`values.yaml`

```
defaultBackend:  
  name: default-backend  
  port: 3000  
  image:  
    repository: myrepo/backend-image  
    tag: 'v1'
```

Container port

Container image name and tag

```
helm install myingress \  
  stable/nginx-ingress \  
  -f values.yaml \  
  -n kube-ingress
```



ingress-nginx Features

```
apiVersion: networking.k8s.io/v1
kind: Ingress
metadata:
  name: search-ingress
  annotations:
    nginx.ingress.kubernetes.io/canary: "true"
    nginx.ingress.kubernetes.io/canary-weight: "20"
```

```
spec:
  ingressClassName: nginx
  rules:
  - host: acme.com
    http:
      paths:
      - path: /search
        pathType: Prefix
        backend:
          name: search-v2
          port:
            number: 8080
```

Redirect 20% of the traffic to /search



nginx-ingress Features

- Enable CORS to response

```
nginx.ingress.kubernetes.io/enable-cors: "true"
```

- Rate limit the number of request from a given IP per minute and seconds. Returns a 503 if threshold is breached

```
nginx.ingress.kubernetes.io/limit-rps: "5"  
nginx.ingress.kubernetes.io/limit-rpm: "300"
```

- Enable affinity/stickiness

```
nginx.ingress.kubernetes.io/affinity: "cookie"  
nginx.ingress.kubernetes.io/affinity-mode: "persistent"  
nginx.ingress.kubernetes.io/session-cookie-name: "sessionid"
```



ingress-nginx Features

- Deploy ingress-nginx with modsecurity module enabled
 - Modsecurity enables web application firewall (L7 firewall)

values.yaml

```
controller:
  config:
    enable-modsecurity: "true"
    enable-owasp-modsecurity-crs: "true"
```

```
helm install my-ingress \
  stable/nginx-ingress -f values.yaml
```

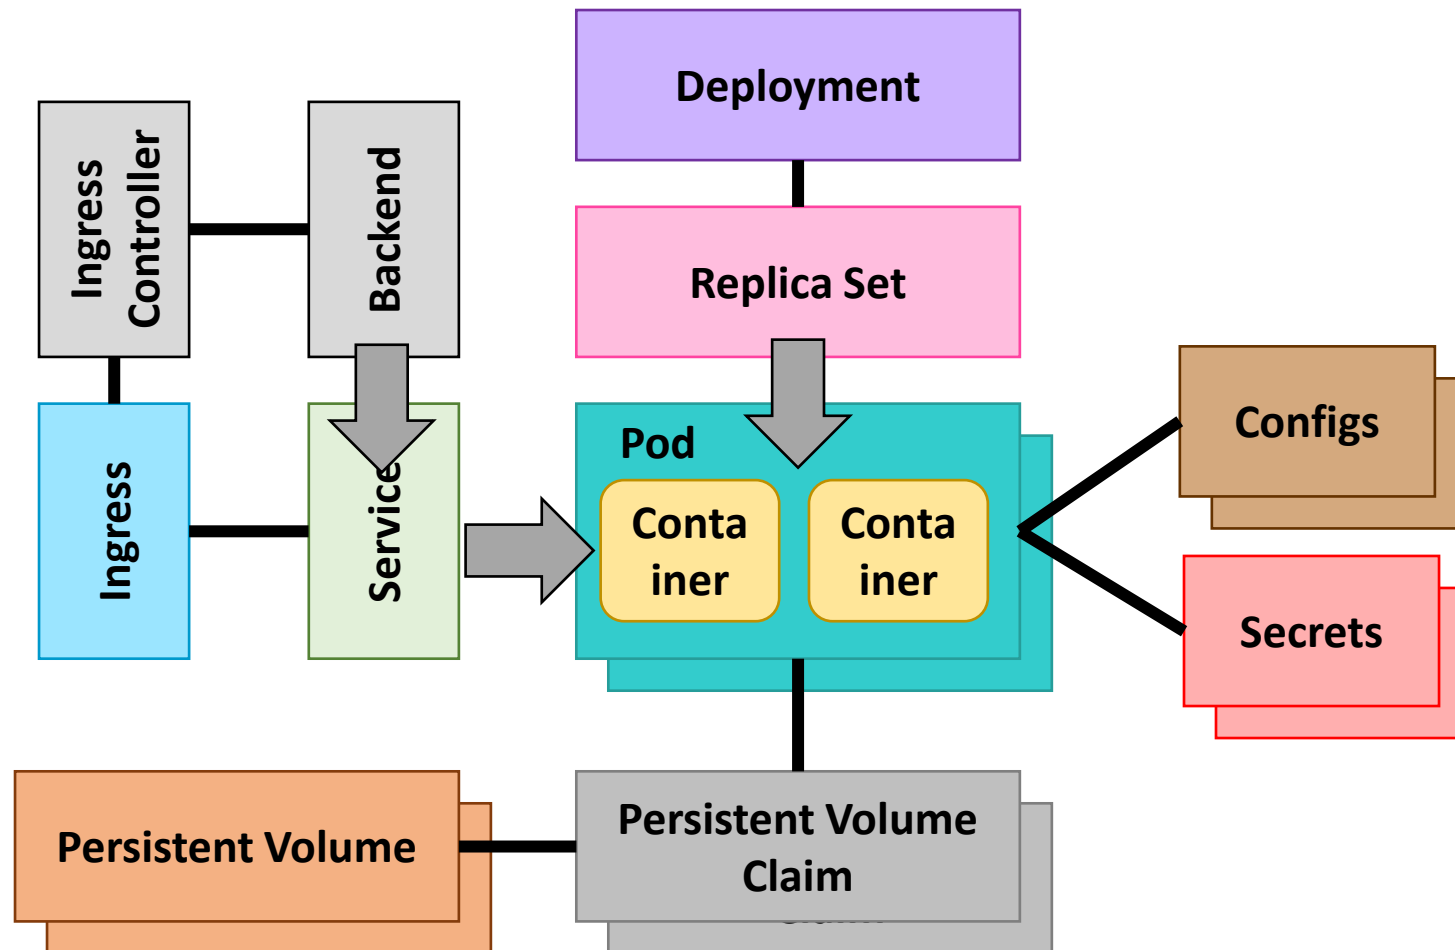
- WAF is enable for all routes by default
 - Need to explicitly disable it
 - See <https://kubernetes.github.io/ingress-nginx/user-guide/nginx-configuration/annotations/#modsecurity>

```
nginx.ingress.kubernetes.io/enable-modsecurity: "false"
```

```
nginx.ingress.kubernetes.io/enable-owasp-core-rules: "false"
```

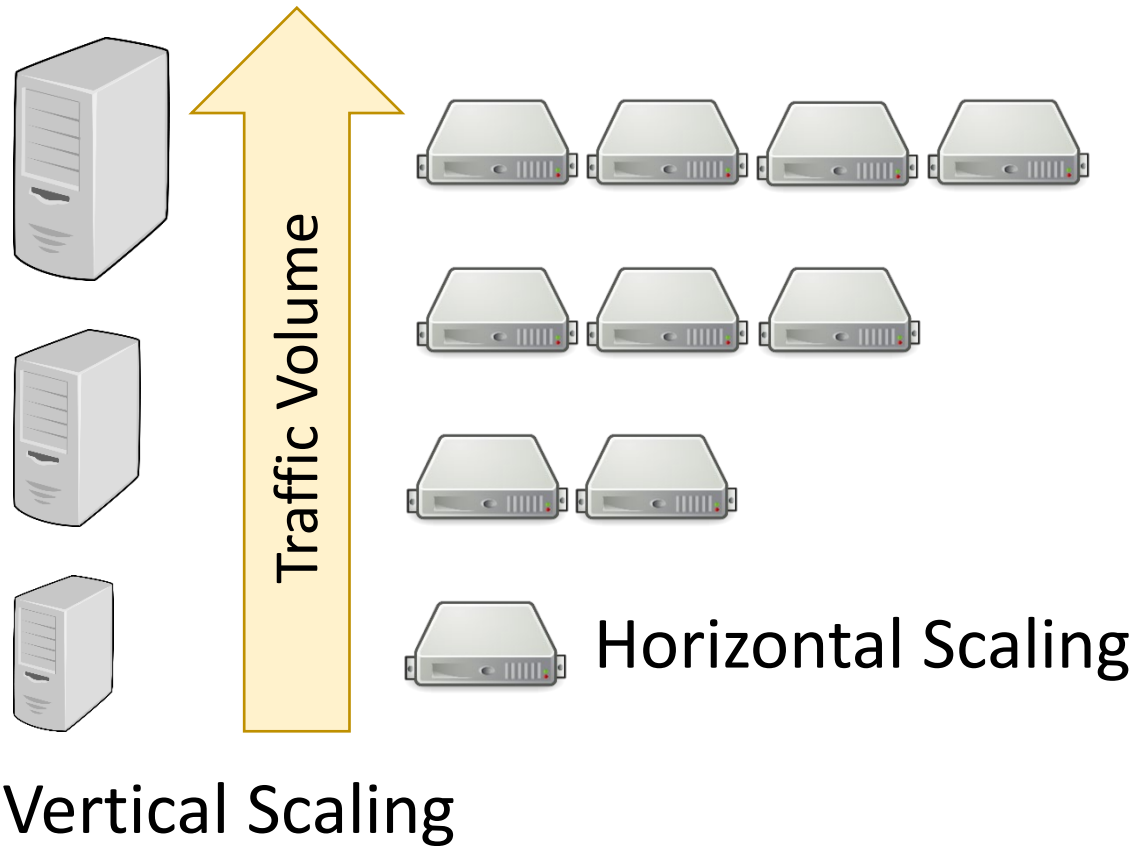


Kubernetes





Scaling



- Scaling is the capability of the system to handle more workload by provisioning more resources
- Three types of scaling
 - Horizontal scaling - scales by provision more Pods
 - Applications must be stateless allowing the ingress controller to route the request to any Pod
 - Vertical scaling - scaling by migrating the application to a 'larger' node
 - Application must be able to utilize the extra resources eg. more vCPUs or memory
 - Cluster scaling - scale cluster by adding nodes to the cluster
 - Cloud provider specific - typically configured when provisioning the cluster



Why Scale?

- Efficient use of resources
 - Ensure that the actual usage is on parity with the current usage
- Dynamically respond to workload fluctuation
 - Elasticity - providing an agreed on SLA
- Cost optimization
 - Pay only what you use



Horizontal Manual Scaling

- Types of scaling
 - Manual
 - Automatic - Horizontal Pod Autoscaler
- Use `kubectl` to scale up or down

```
kubectl scale --replicas <number> deployment <deployment>
```



Metrics Server

- Need to collect metrics for Kubernetes to make decision on scaling
- metrics-server is a set of pods running in Kubernetes
 - Collects CPU and memory utilization
 - Stores them in memory not to an external datastore
 - For viewing historical data, require more advance packages like Prometheus
- Can be installed from a YAML file or as a helm chart



Monitoring Kubernetes with top

- Node metrics

```
kubectl top node
```

- Pod metrics

```
kubectl top pod
```

- `top` will only work if metrics-server has been installed



Horizontal Pod Autoscaler

- HPA scales a deployment based on one or more metrics
 - Eg. trigger scaling when CPU utilization breaches 80%
 - Metrics to scale the Pods can be
 - Build in metrics , custom metrics, external metrics
- HPA runs a control loop - runs every 30 seconds (default)
 - Queries metrics server
 - Match that against the specified threshold
 - Updates the number of replicas in a deployment if required to meet the load
 - Deployment would then perform the scaling (in or out)
- Reduces cluster size if utilization is low for a period of time
 - Scaling in

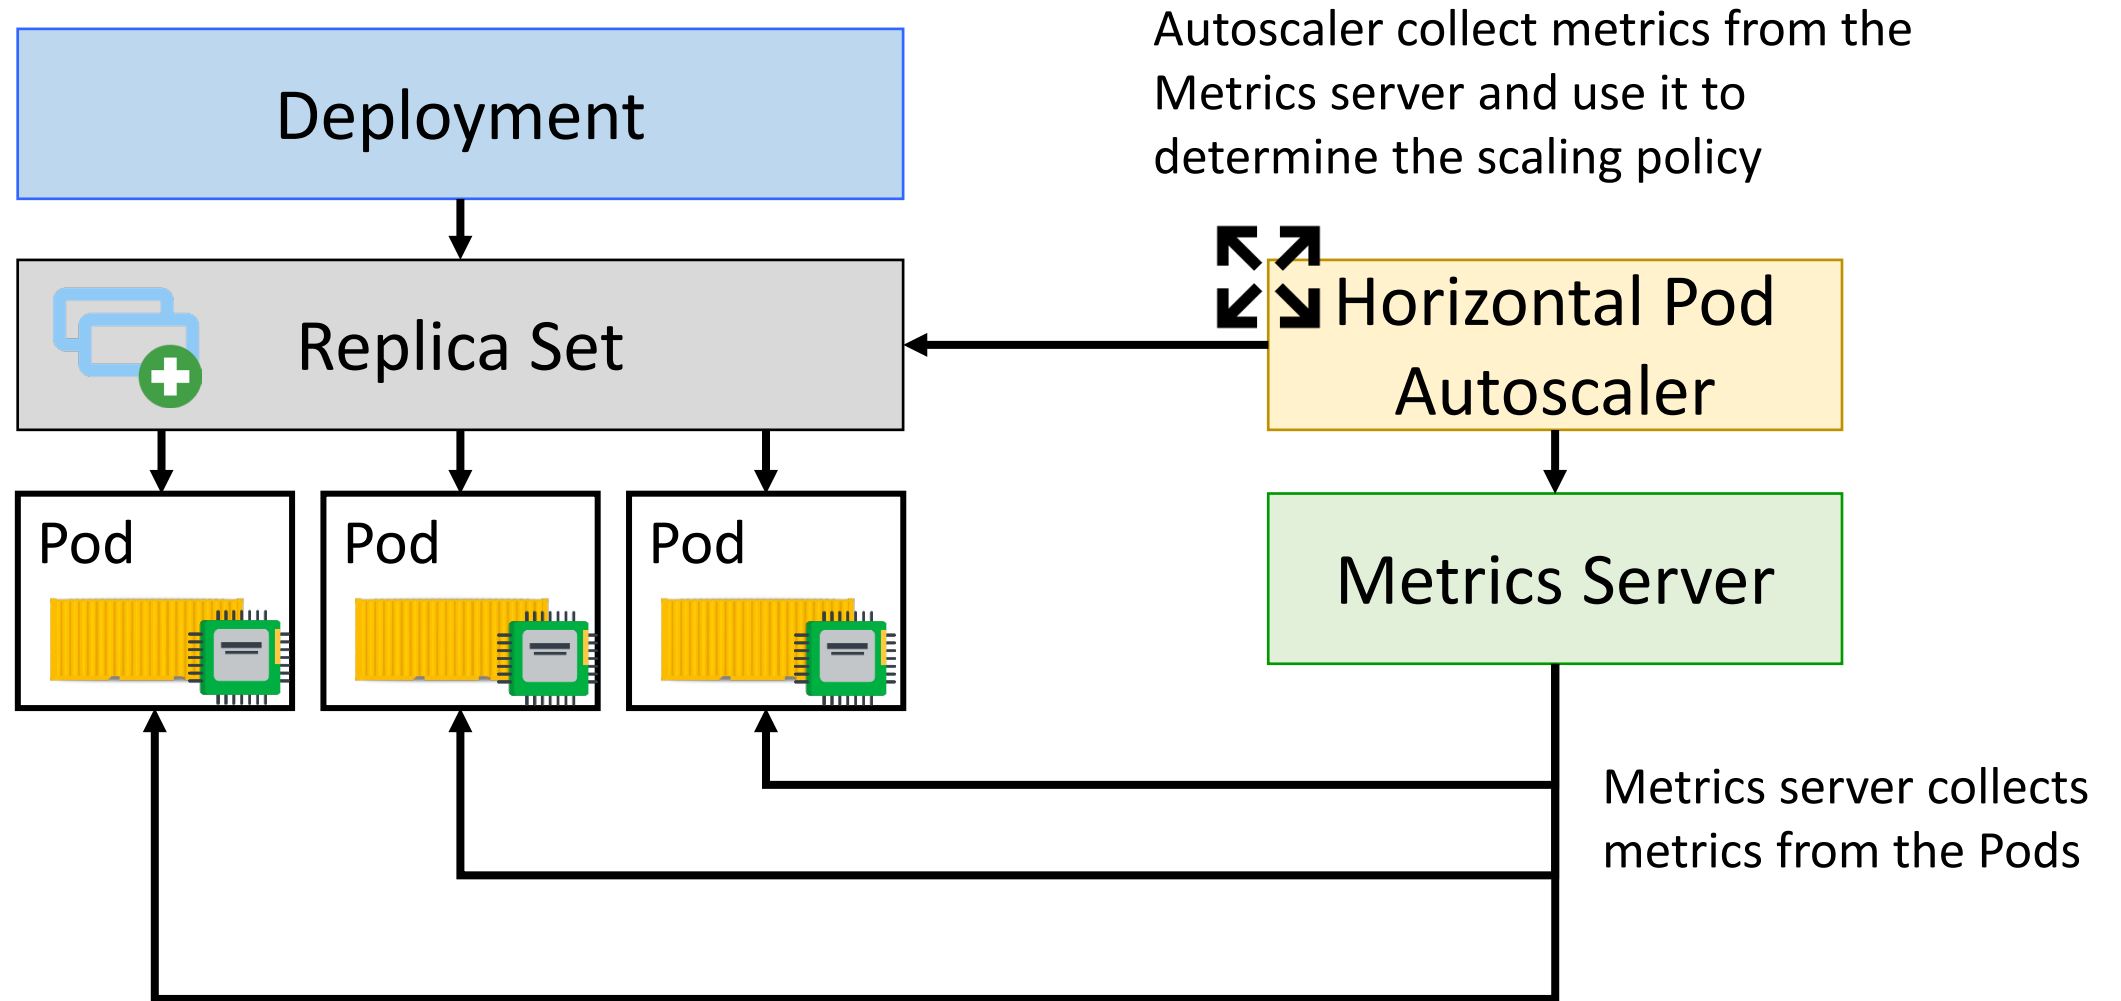


Setting Pod Request

- Horizontal scaler scales a Pod by determining if the Pod has breached a certain threshold
 - For memory and CPU
- Set the request for CPU and memory
 - Specify the minimum amount of compute resources required
- Resource type
 - CPU - measured in CPU units eg 100m is 100 millicores
 - 1 CPU in Kubernetes == 1 vCPU, Core, vCore, Hyperthread
 - Memory - 16M



Horizontal Pod Autoscaler





Requesting Resources

```
apiVersion: v1
kind: Pod
metadata:
  name: myapp
spec:
  containers:
  - name: myapp
    image: myapp:sha256:...
```

HPA only looks
at the CPU

```
resources:
```

```
requests:
```

```
  cpu: 100m
```

```
  memory: 16M
```

```
limits:
```

```
  memory: 32M
```

```
...
```

} Request the minimum amount of
compute resources

→ Describe the maximum amount
of compute resources required



Defining a Horizontal Pod Autoscaler

```
apiVersion: autoscaling/v2beta2
kind: HorizontalPodAutoScaler
metadata:
  name: myapp
spec:
  minReplicas: 1
  maxReplicas: 8
  scaleTargetRef:
    apiVersion: apps/v1
    kind: Deployment
    name: myapp
```

The deployment that
this HPA is targeting

Minimum and maximum number of replicas. Since the HPA is managing the replica set, this setting takes precedence over the deployment setting

metrics:

- type: Resource
resource:
name: cpu
target:
type: Utilization
averageUtilization: 80
- type: Resource
resource:
name: memory
target:
type: Utilization
averageUtilization: 80

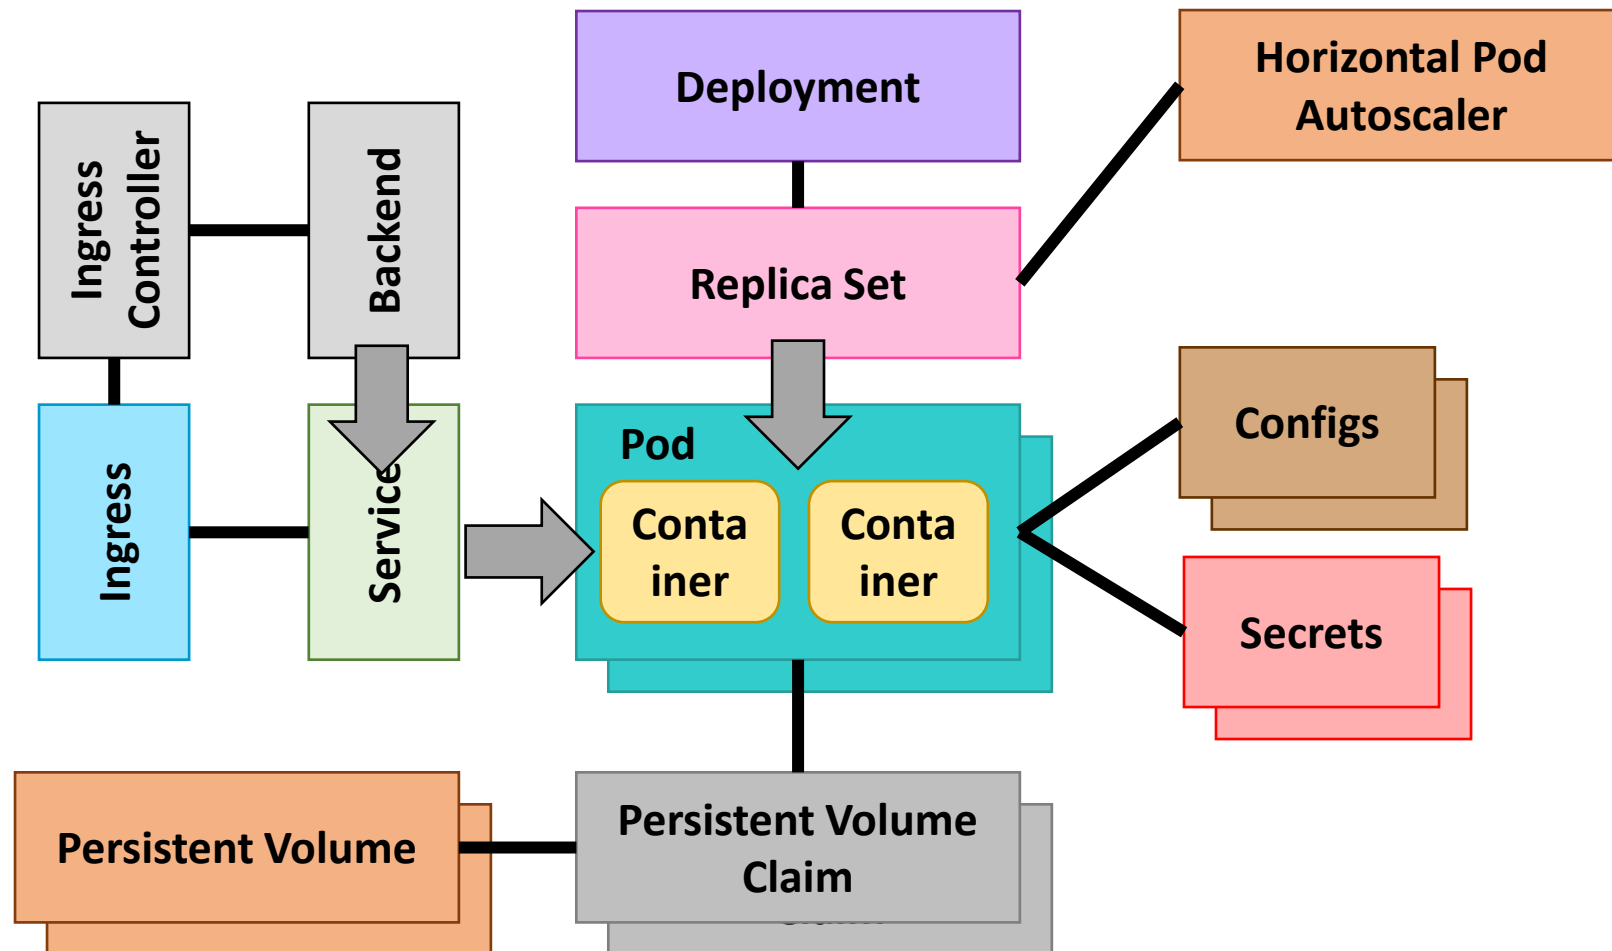
Scale the pod
when these
threshold are
breached

80

80



Kubernetes





Appendix



Using ConfigMaps


Injecting as environment variables

```
containers:
  ...
  env:
    - name: DB_NAME
      valueFrom:
        configMapKeyRef:
          name: myapp-config
          key: db_name
    - name: DB_HOST:
      valueFrom:
        configMapKeyRef:
          name: myapp-config
          key: db_host
```

Mounting as a volume

```
volumes:
  - name: config-volume
    configMap:
      name: myapp-config

containers:
  ...
  volumeMounts:
    - name: config-config
      mountPath: /config
```





Managing Context

- For grouping access parameters under a common name
 - Like a profile
 - Set the namespace, do not need the `-n` option

- Create a context

```
kubectl config set-context <context_name> --namespace=<name> \
    --cluster=<cluster_name> --user=<user_name>
```

- View current contexts

```
kubectl config view
```

- Use a context

```
kubectl config use-context <context_name>
```