




Kubernetes

Part 2



Kubernetes





Volumes

- Volumes are storage that are shared by containers in a Pod
 - Allocated by the Pod, usually a shared directory in the Pod
 - Not visible outside of Pod
- Tied to the lifecycle of a Pod viz. its removed when the Pod is delete
 - Unlike Docker volumes where they are durable
- Different types of volumes
 - Eg. hostPath, NFS, iSCSI, fibre channel, empty directory, etc.
- **hostPath** and **emptyDir** type is good for sharing data between containers in a Pod
 - Eg. The example of file puller and web server


emptyDir will be deleted when pod is deleted



hostPath is similar to the concept of docker bind mount

Defining a Volume

Same syntax for creating for Pod templates in deployment objects





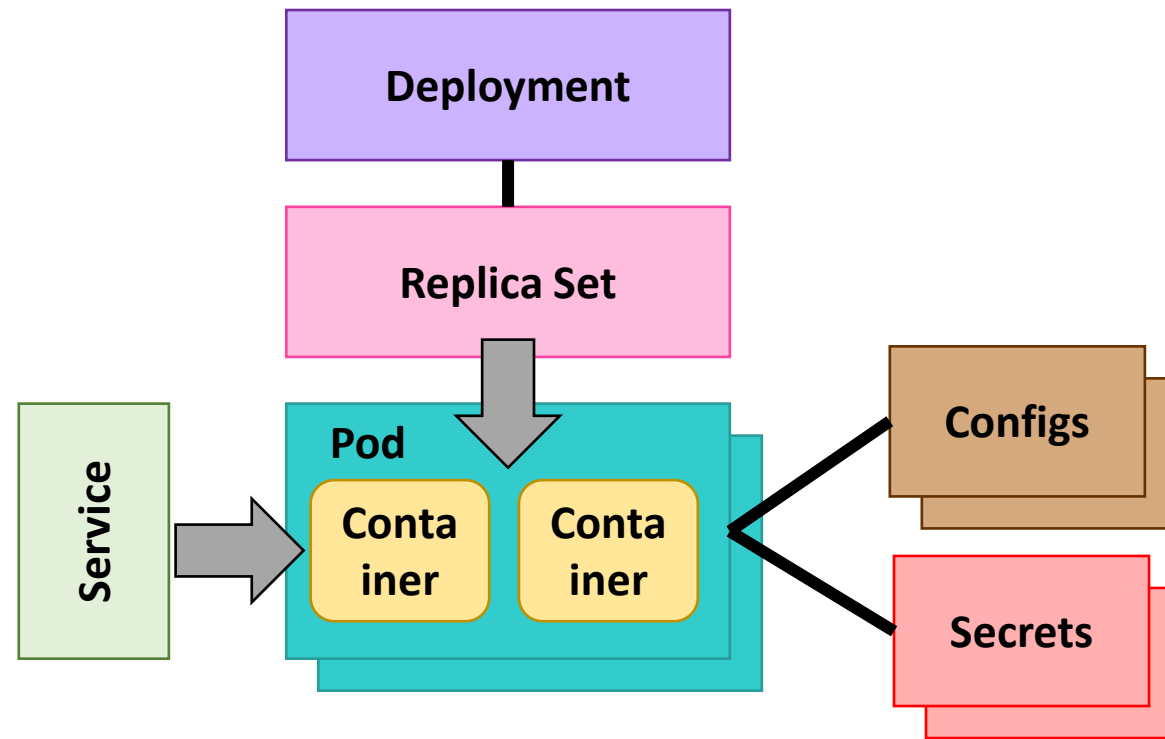
Mounting ConfigMaps and Secrets

- `configMap` and `secrets` can be mounted as volumes
 - Keys becomes the filename, the value is the file content
- Use cases
 - Passing configuration files
eg `nginx.conf`

```
spec:
  volumes:
    - name: html-vol
      configMap:
        name: html-assets
  containers:
    - name: nginx
      image: nginx
      volumeMounts:
        - name: html-vol
          mountPath: /usr/share/nginx/html
          readOnly: true
```



Kubernetes





Persistent Storage

- Kubernetes can dynamically provision storage
 - Eg. User ask for 50GB volume to caching images
- Kubernetes allows storage to be either statically or dynamically provisioned
 - Static provision - an administrator will need to first provision the storage manually
 - Dynamic provision - the user describes the type of storage that is required; Kubernetes will attempt to provision based on the user's requirements
- Once a persistent storage has been allocated and claimed/reserved, a Pod can mount the volume like any regular volume
- Persistent volumes lifecycle are not tied to the Pod's lifecycle
 - Unlike volumes, persistent volumes will not be deleted when a Pod is deleted
 - This behaviour can be configured

Persistent volume claims will not be deleted when pod is deleted



Key Concepts

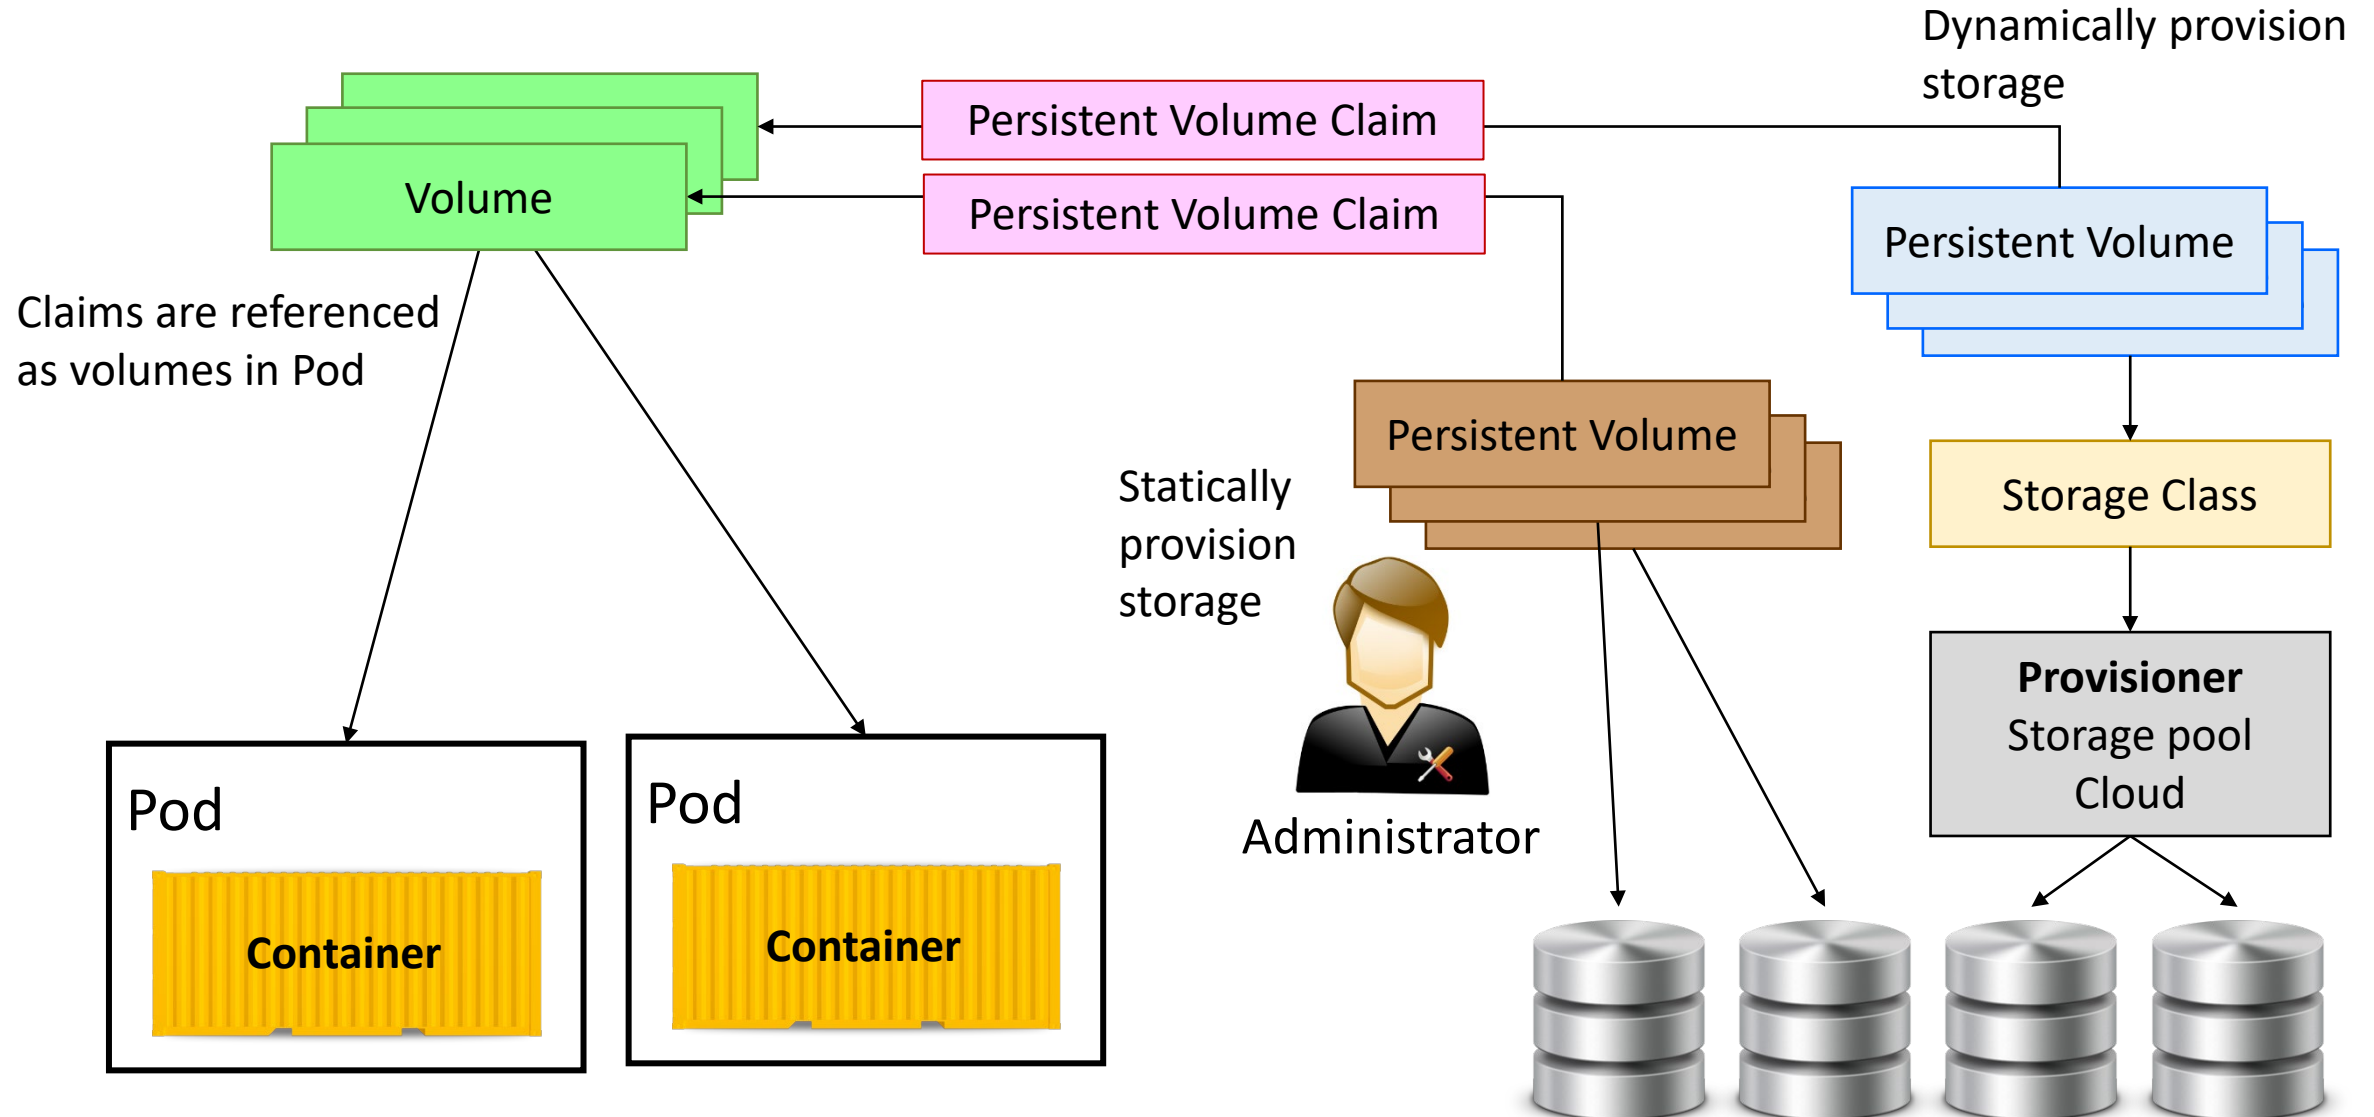
- Storage class - a type of storage
 - Who the provisioner, storage specific details, retention policy, etc.
- Persistent volume - the actual storage
 - A piece of storage provisioned by an administrator or thru storage class
 - Supports may different storage type
 - AWS EBS, Azure File Service, Cinder, fibre channel, GCP Disk, NFS, etc.
 - Different type of access mode - exclusive or shared
- Persistent Volume claim - when a persistent volume has been allocated for use, the volume is staid to be claimed



If ReclaimPolicy of StorageClass is 'Delete', when persistent volume claim is deleted, volume provisioned will also be deleted

Persistent Volume

If ReclaimPolicy of StorageClass is 'Retain', when persistent volume claim is deleted, volume will be retained





Static vs Dynamic

Static

- Administrator has to manually allocate storage and map it to a persistent volume
- Users can then claim this volume

Dynamic

- When Kubernetes tries to resolve a claim and the persistent volume is unavailable
- It looks for a storage class that best matches the request storage
- Dynamically creates the persistent volume using the provisioner



Defining a Persistent Volume Claim

```
apiVersion: v1
```

```
kind: PersistentVolumeClaim
```

```
meta-data:
```

```
  name: myapp-pvc
```

```
spec:
```

```
  accessModes:
```

```
    - ReadWriteOnce
```

List of access modes

ReadOnlyMany

ReadWriteMany

```
resources:
```

```
  requests:
```

```
    storage: 5Gi
```

```
storageClassName: standard
```

Get storage class name(s) with
kubectl get storageclass



Mounting a Persistent Volume


```
apiVersion: v1
kind: Pod
meta-data:
  name: myapp
```

```
spec:
```

```
  volumes:
```

```
    - name: data-volume
      persistentVolumeClaim:
        claimName: myapp-pvc
```

Specify the claim name



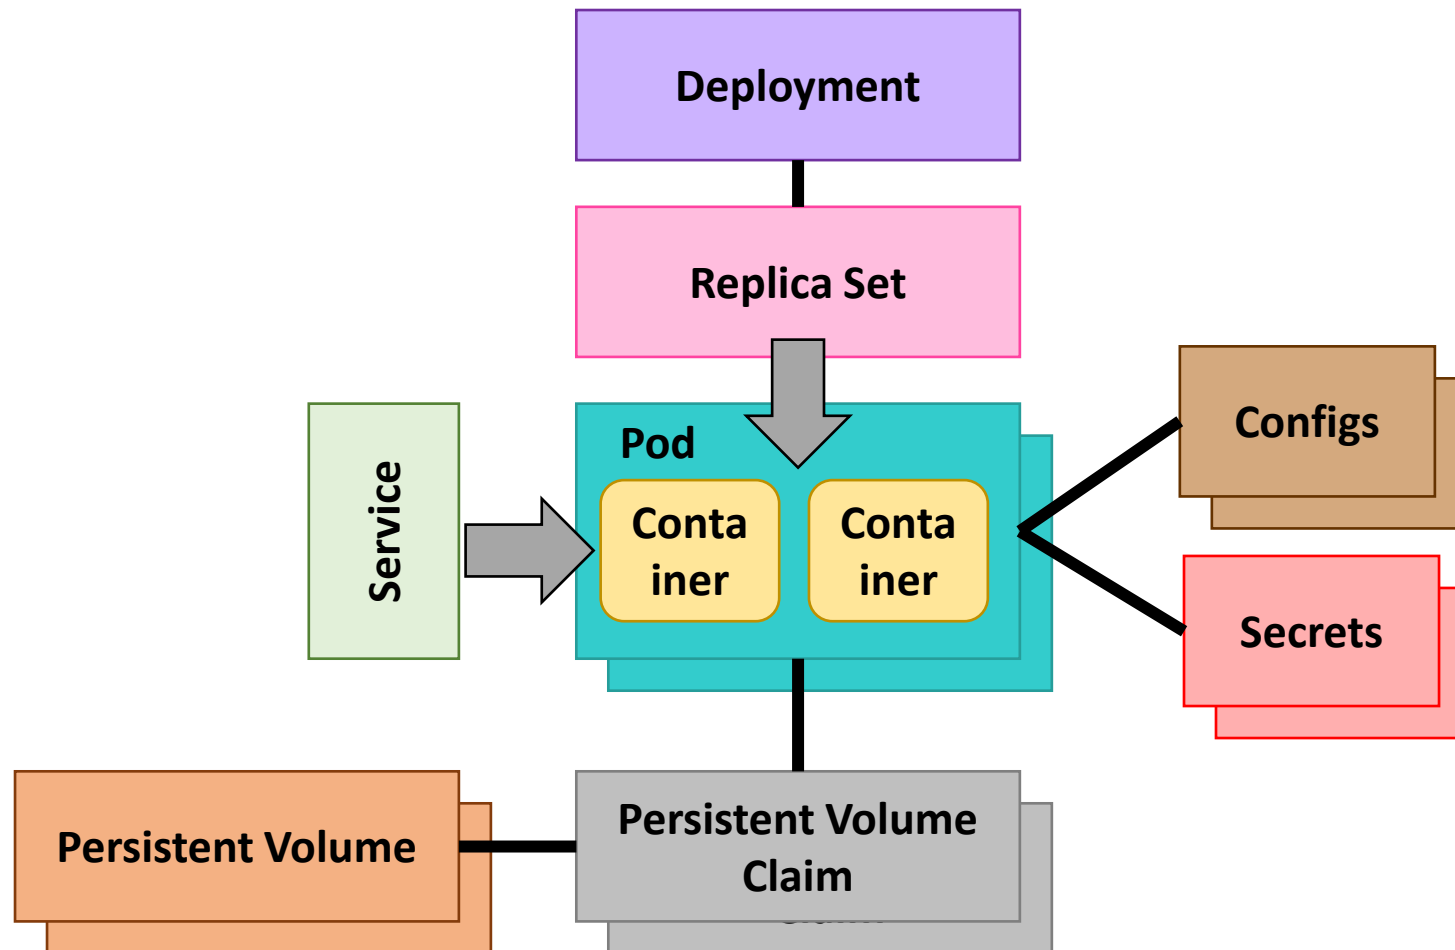
```
  containers:
```

```
    - name: myapp
      volumeMounts:
        - mountPath: /app/public
          name: data-volume
```

```
    ...
```



Kubernetes





Persistence Volume Management

- Display persistence volume detail
 - Persistence volume - `kubectl get pv`
 - Persistence volume claim - `kubectl get pvc`
 - Storage classes - `kubectl get sc`
- Delete persistence volume


```
kubectl delete pvc <name>
```

```
kubectl delete pv <name>
```



Load Balancer and Ingress

- By default services are allocated a cluster IP
 - Only accessible within the cluster
- Load balancer exposes the service to the public
 - Accessible from outside of the cluster
 - Load balancer will redirect the request to pods based on its routing policy
 - Another way to allow external access is via node port
- Load balancer are resources that are provisioned from the underlying cloud platform
 - May have more features that you require
 - Also cost more





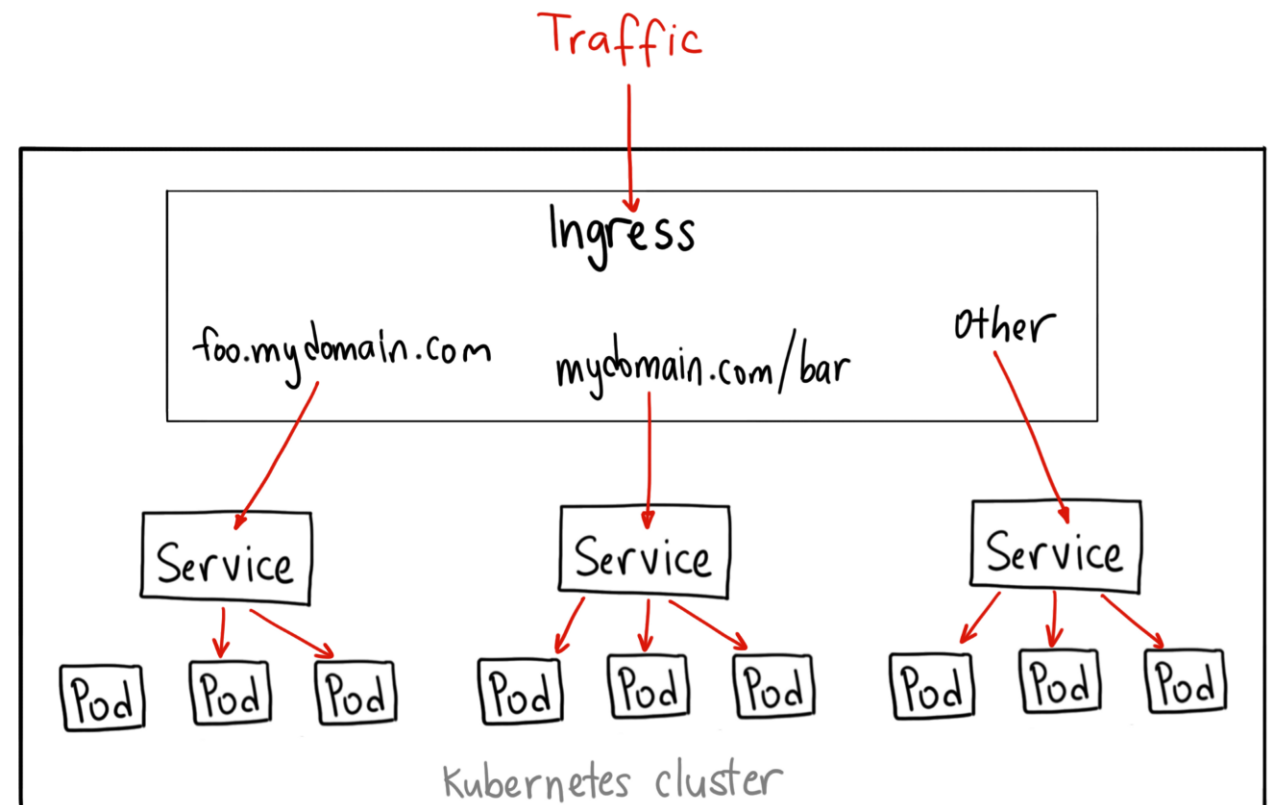
Load Balancer and Ingress

- An ingress, like a LoadBalancer service, allows traffic into the cluster
 - Provisions an 'external' load balancer
 - Shared by many services within the cluster
 - More cost effective
- May require to provision an Ingress controller
 - May not be available
- NGINX Ingress controller is a popular ingress controller
 - Deploys NGINX as Ingress
 - <https://github.com/kubernetes/ingress-nginx>



Ingress

- Application layer (L7) router that sits in front of multiple services
- Define a set of routing rules on how services are access externally
 - Eg. 2 services, one for search one for checkout. Might map to /search and /checkout
- Rules are applied to ingress controllers which performs the actual routing
 - Controllers might be a cloud provider's load balancer or Nginx reverse-proxy





Defining an Ingress

```
apiVersion: networking.k8s.io/v1
```

```
kind: Ingress
```

```
metadata:
```

```
  name: myapp
```

```
  annotations:
```

```
    nginx.ingress.kubernetes.io/rewrite-target: "/"
```

```
    nginx.ingress.kubernetes.io/ssl-redirect: "false"
```

```
spec:
```

```
  ingressClassName: nginx
```

```
  rules:
```

```
    - host: acme.com
```

```
      http:
```

```
        paths:
```

```
          - path: /hello
```

```
            pathType: Prefix
```

```
            backend:
```

```
              service:
```

```
                name: myapp
```

```
                port:
```

```
                  number: 8080
```

Change/rewrite a matched resource
to its root e.g /hello to /

Used to configure
NGINX ingress
controller

One or more of
these rules to
specify which
services to handle
what resource

Select the ingress controller to use

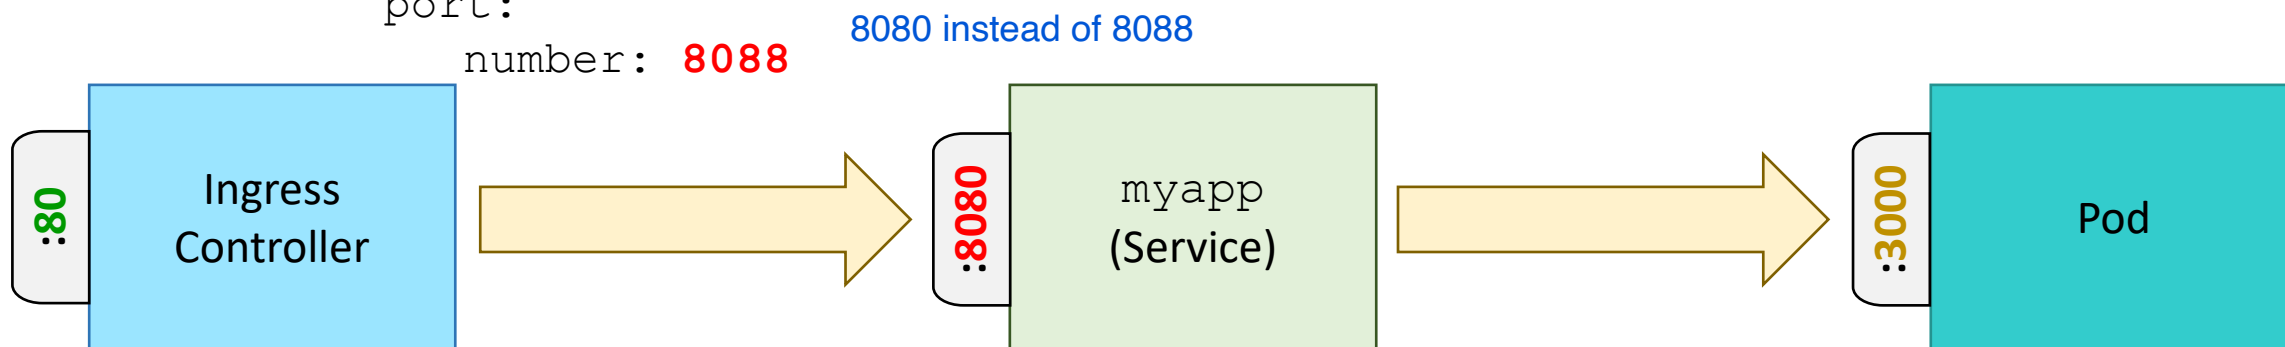


Ingress Ports

```
kind: Ingress
spec:
  rules:
  - host: acme.com
    http:
      paths:
      - path: /
        pathType: Prefix
        backend:
          service:
            name: mysvc
            port:
              number: 8088
```

```
kind: Service
spec:
  ports:
  - port: 8080
    targetPort: 3000
```

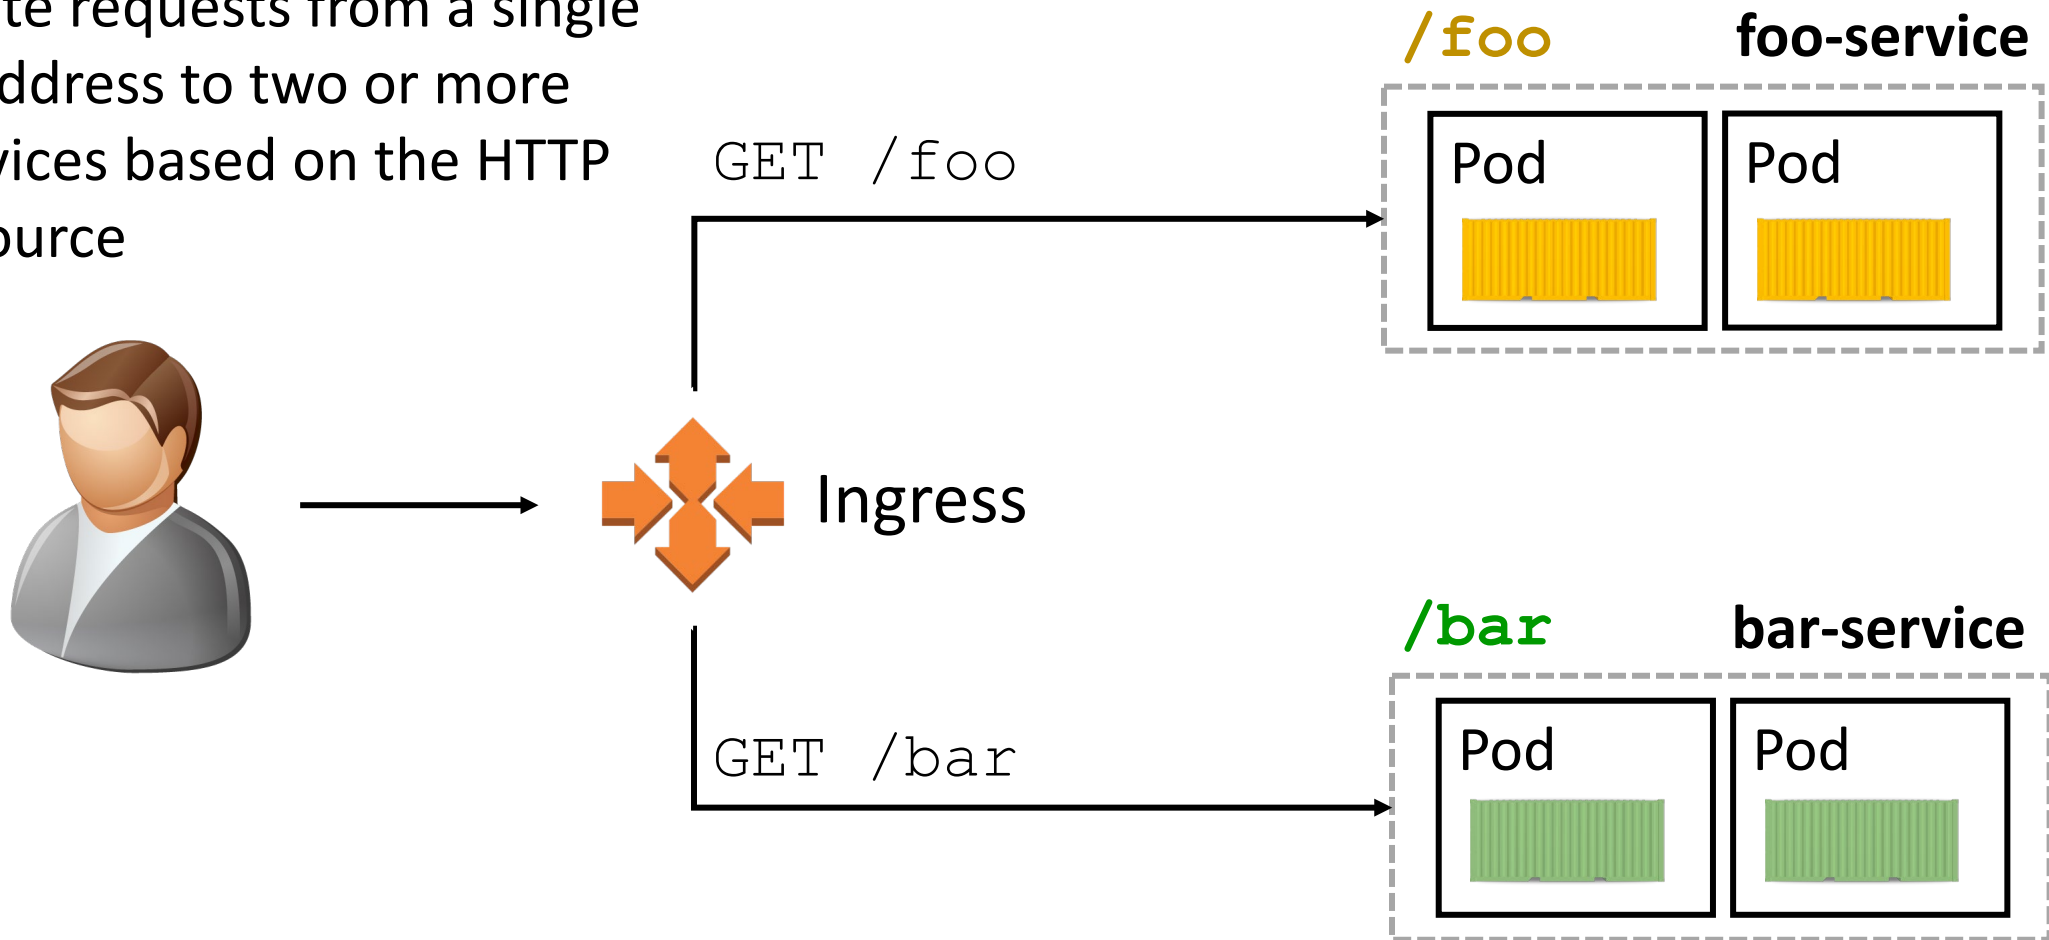
```
kind: Deployment
spec:
  containers:
  - containerPort: 3000
```





Ingress - Fan Out

Route requests from a single IP address to two or more services based on the HTTP resource





Ingress Fan Out Example

```
apiVersion: networking.k8s.io/v1
kind: Ingress
metadata:
  name: myapp-ingress
  annotations:
    nginx.ingress.kubernetes.io/rewrite-target: /
spec:
  ingressClassName: nginx
  rules:
  - host: acme.com
    http:
      paths:
      - path: /foo
        pathType: Prefix
        backend:
          service:
            name: foo-service
            port:
              number: 8000
      - path: /bar
        pathType: Prefix
        backend:
          service:
            name: bar-service
            port:
              number: 8001
```

Note: If the 2 services are web application, then all resources references (eg in HTML) are now rooted under `/foo` or `/bar`

So the web application must take this into account

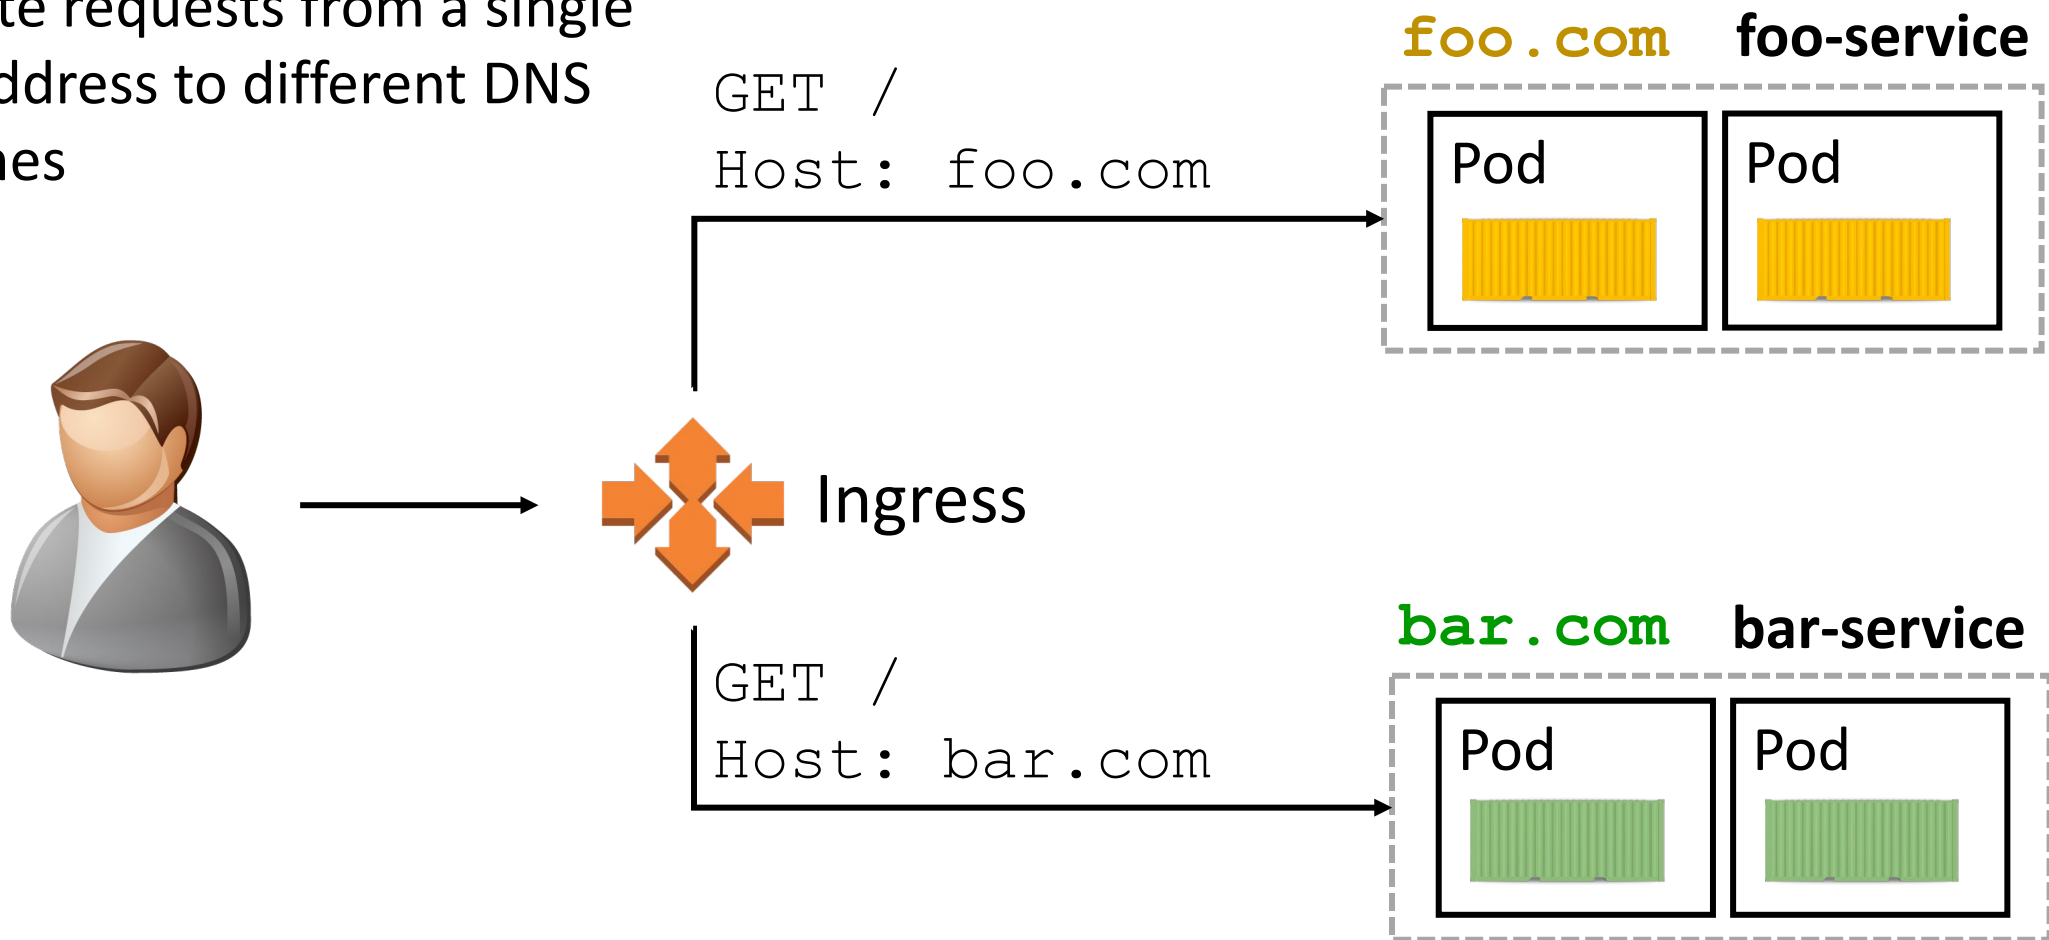
One option is to use relative reference or use `<base>`

Request is routed to these 2 services
`acme.com/foo` or `acme.com/bar`



Ingress - Virtual Host

Route requests from a single IP address to different DNS names





Ingress Virtual Host Example

```
apiVersion: networking.k8s.io/v1
```

```
kind: Ingress
```

```
metadata:
```

```
  name: myapp-ingress
```

```
spec:
```

```
  ingressClassName: nginx
```

```
  rules:
```

```
    - host: foo.com
```

```
      http:
```

```
        paths:
```

```
          - pathType: Prefix
```

```
            backend:
```

```
              name: foo-service
```

```
              port:
```

```
                number: 8080
```

```
    - host: bar.com
```

```
      http:
```

```
        paths:
```

```
          - pathType: Prefix
```

```
            backend:
```

```
              name: bar-service
```

```
              port:
```

```
                number: 8080
```

foo.com

host

bar.com

host

Request is routed to these 2 services depending on the `Host` attribute.



Kubernetes Annotations

- Annotations are additional/proprietary metadata passed to a controller
 - Usually configuration information
- Allow additional configuration not part of the spec:
 - Eg. ingress resources does not support canary deployment. Can configure canary if using ingress-nginx thru annotations
- <https://kubernetes.github.io/ingress-nginx/user-guide/nginx-configuration/annotations>



ingress-nginx Features

```
apiVersion: networking.k8s.io/v1
```

```
kind: Ingress
```

```
metadata:
```

```
  name: myapp-ingress
```

```
  annotations:
```

```
    nginx.ingress.kubernetes.io/rewrite: /$2
```

```
spec:
```

```
  ingressClassName: nginx
```

```
  rules:
```

```
    - host: acme.com
```

```
      http:
```

```
        paths:
```

```
          - path: /foo(/|$) (.*)
```

```
            pathType: Prefix
```

```
            backend:
```

```
              ...
```

```
          - path: /bar(/|$) (.*)
```

```
            pathType: Prefix
```

```
            backend:
```

Remove the first segment of the resource from the request

Original URL	Rewrote URL
/foo	/
/foo/debug	/debug
/foo/api/customer/123	/api/customer/123

Original URL: /foo/abc

New URL: /abc

\$2 refers to characters after /foo/ (e.g. abc)



Handling Errors

- If no rules matches the incoming request, then the traffic can be routed to a default backend if it is configured
 - The default backend service can be any of your application
 - Eg. Help page, chatbot

Traffic will be routed to this service if no rule matches
If the default backend is not specified, then will fallback to the ingress controller

```
apiVersion: v1
kind: Ingress
metadata:
  name: myapp-ingress
  annotations:
    nginx.ingress.kubernetes.io/rewrite: /$2
spec:
  defaultBackend:
    service:
      name: defaultBackendService
      port:
        number: 8080
  rules:
  - host: acme.com
    http:
      ...
```

A curved arrow originates from the text 'Traffic will be routed to this service if no rule matches' and points directly to the 'defaultBackend' field in the YAML code block.



Handling Errors

- An alternative is to install a global default service
 - This feature is specific to the ingress controller
- For `stable/nginx-ingress`, this is done during deployment
 - With helm
- Can only have a single default backend

values.yaml

```
defaultBackend:  
  name: default-backend  
  port: 3000  
  image:  
    repository: myrepo/backend-image  
    tag: 'v1'
```

Container port

Container image name and tag

```
helm install myingress \  
  stable/nginx-ingress \  
  -f values.yaml \  
  -n kube-ingress
```



ingress-nginx Features

```
apiVersion: networking.k8s.io/v1
kind: Ingress
metadata:
  name: search-ingress
  annotations:
    nginx.ingress.kubernetes.io/canary: "true"
    nginx.ingress.kubernetes.io/canary-weight: "20"
```

```
spec:
  ingressClassName: nginx
  rules:
  - host: acme.com
    http:
      paths:
      - path: /search
        pathType: Prefix
        backend:
          name: search-v2
          port:
            number: 8080
```

Redirect 20% of the traffic to /search

Spin up two ingress resources
A main ingress and a canary ingress
Canary ingress will handle 20% of traffic



nginx-ingress Features

- Enable CORS to response

```
nginx.ingress.kubernetes.io/enable-cors: "true"
```

- Rate limit the number of request from a given IP per minute and seconds. Returns a 503 if threshold is breached

```
nginx.ingress.kubernetes.io/limit-rps: "5"  
nginx.ingress.kubernetes.io/limit-rpm: "300"
```

- Enable affinity/stickiness

```
nginx.ingress.kubernetes.io/affinity: "cookie"  
nginx.ingress.kubernetes.io/affinity-mode: "persistent"  
nginx.ingress.kubernetes.io/session-cookie-name: "sessionid"
```



ingress-nginx Features

- Deploy ingress-nginx with modsecurity module enabled
 - Modsecurity enables web application firewall (L7 firewall)

values.yaml

```
controller:
  config:
    enable-modsecurity: "true"
    enable-owasp-modsecurity-crs: "true"
```

```
helm install my-ingress \
  stable/nginx-ingress -f values.yaml
```

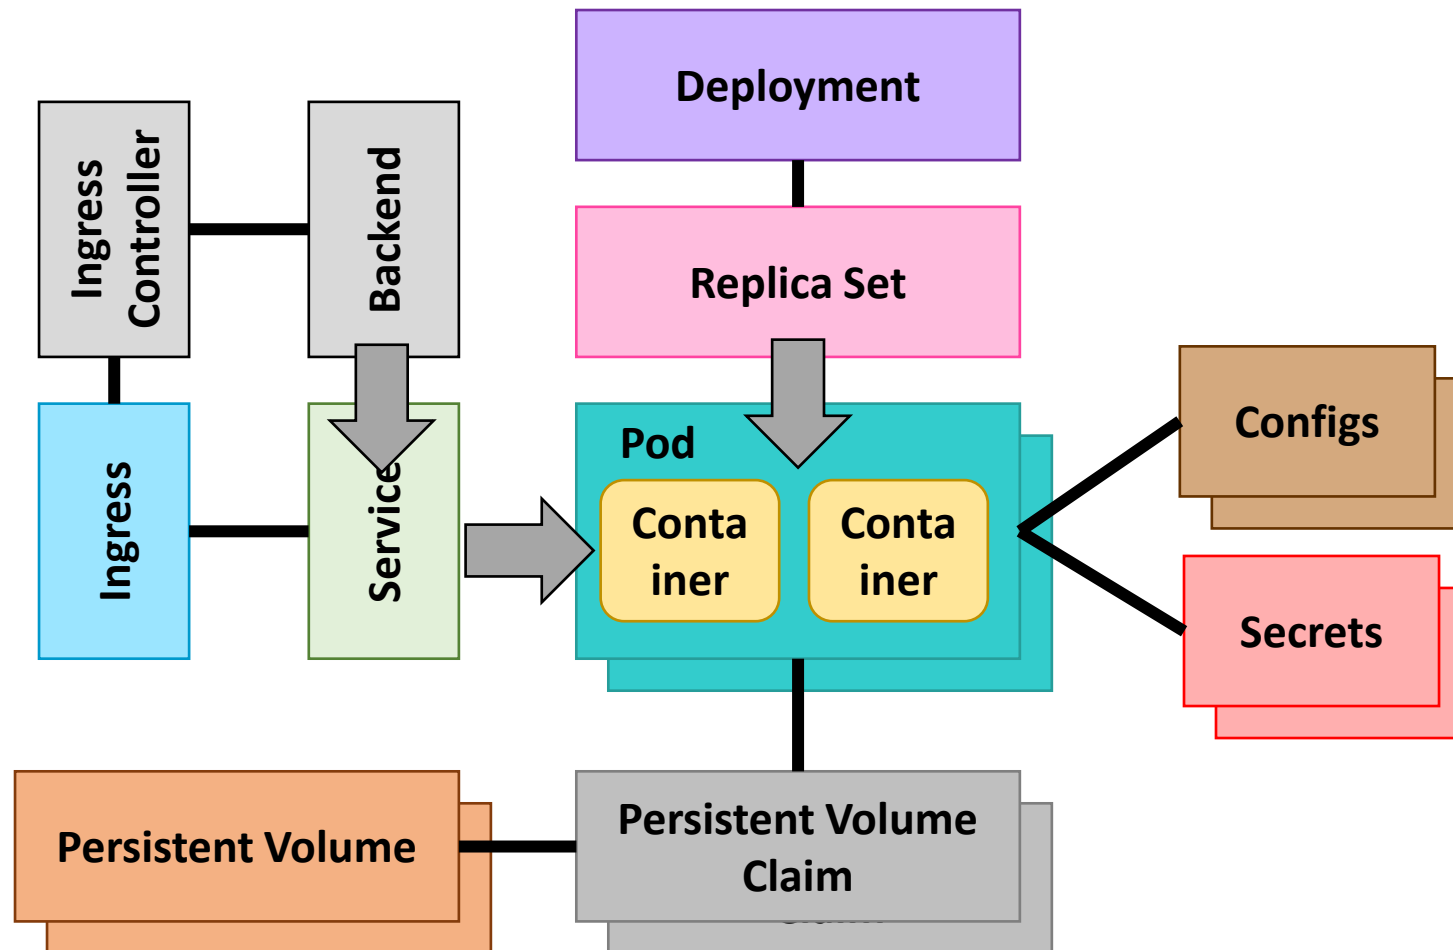
- WAF is enable for all routes by default
 - Need to explicitly disable it
 - See <https://kubernetes.github.io/ingress-nginx/user-guide/nginx-configuration/annotations/#modsecurity>

```
nginx.ingress.kubernetes.io/enable-modsecurity: "false"
nginx.ingress.kubernetes.io/enable-owasp-core-rules: "false"
```



Kubernetes


Ingress controller will be superseded by API gateway
(gateway-api.sigs.k8s.io)





Vertical Pod Autoscaler (VPA) will disrupt service when it tries to migrate application to larger node.

Scaling



- Scaling is the capability of the system to handle more workload by provisioning more resources
- Three types of scaling
 - Horizontal scaling - scales by provision more Pods
 - Applications must be stateless allowing the ingress controller to route the request to any Pod
 - Vertical scaling - scaling by migrating the application to a 'larger' node
 - Application must be able to utilize the extra resources eg. more vCPUs or memory
 - Cluster scaling - scale cluster by adding nodes to the cluster
 - Cloud provider specific - typically configured when provisioning the cluster



Why Scale?

- Efficient use of resources
 - Ensure that the actual usage is on parity with the current usage
- Dynamically respond to workload fluctuation
 - Elasticity - providing an agreed on SLA
- Cost optimization
 - Pay only what you use



Horizontal Manual Scaling

- Types of scaling
 - Manual
 - Automatic - Horizontal Pod Autoscaler
- Use `kubectl` to scale up or down

```
kubectl scale --replicas <number> deployment <deployment>
```



Metrics Server

- Need to collect metrics for Kubernetes to make decision on scaling
- metrics-server is a set of pods running in Kubernetes
 - Collects CPU and memory utilization
 - Stores them in memory not to an external datastore
 - For viewing historical data, require more advance packages like Prometheus
- Can be installed from a YAML file or as a helm chart



Monitoring Kubernetes with top

- Node metrics

```
kubectl top node
```

- Pod metrics

```
kubectl top pod
```

- `top` will only work if metrics-server has been installed



Horizontal Pod Autoscaler

- HPA scales a deployment based on one or more metrics
 - Eg. trigger scaling when CPU utilization breaches 80%
 - Metrics to scale the Pods can be
 - Build in metrics , custom metrics, external metrics
- HPA runs a control loop - runs every 30 seconds (default)
 - Queries metrics server
 - Match that against the specified threshold
 - Updates the number of replicas in a deployment if required to meet the load
 - Deployment would then perform the scaling (in or out)
- Reduces cluster size if utilization is low for a period of time
 - Scaling in

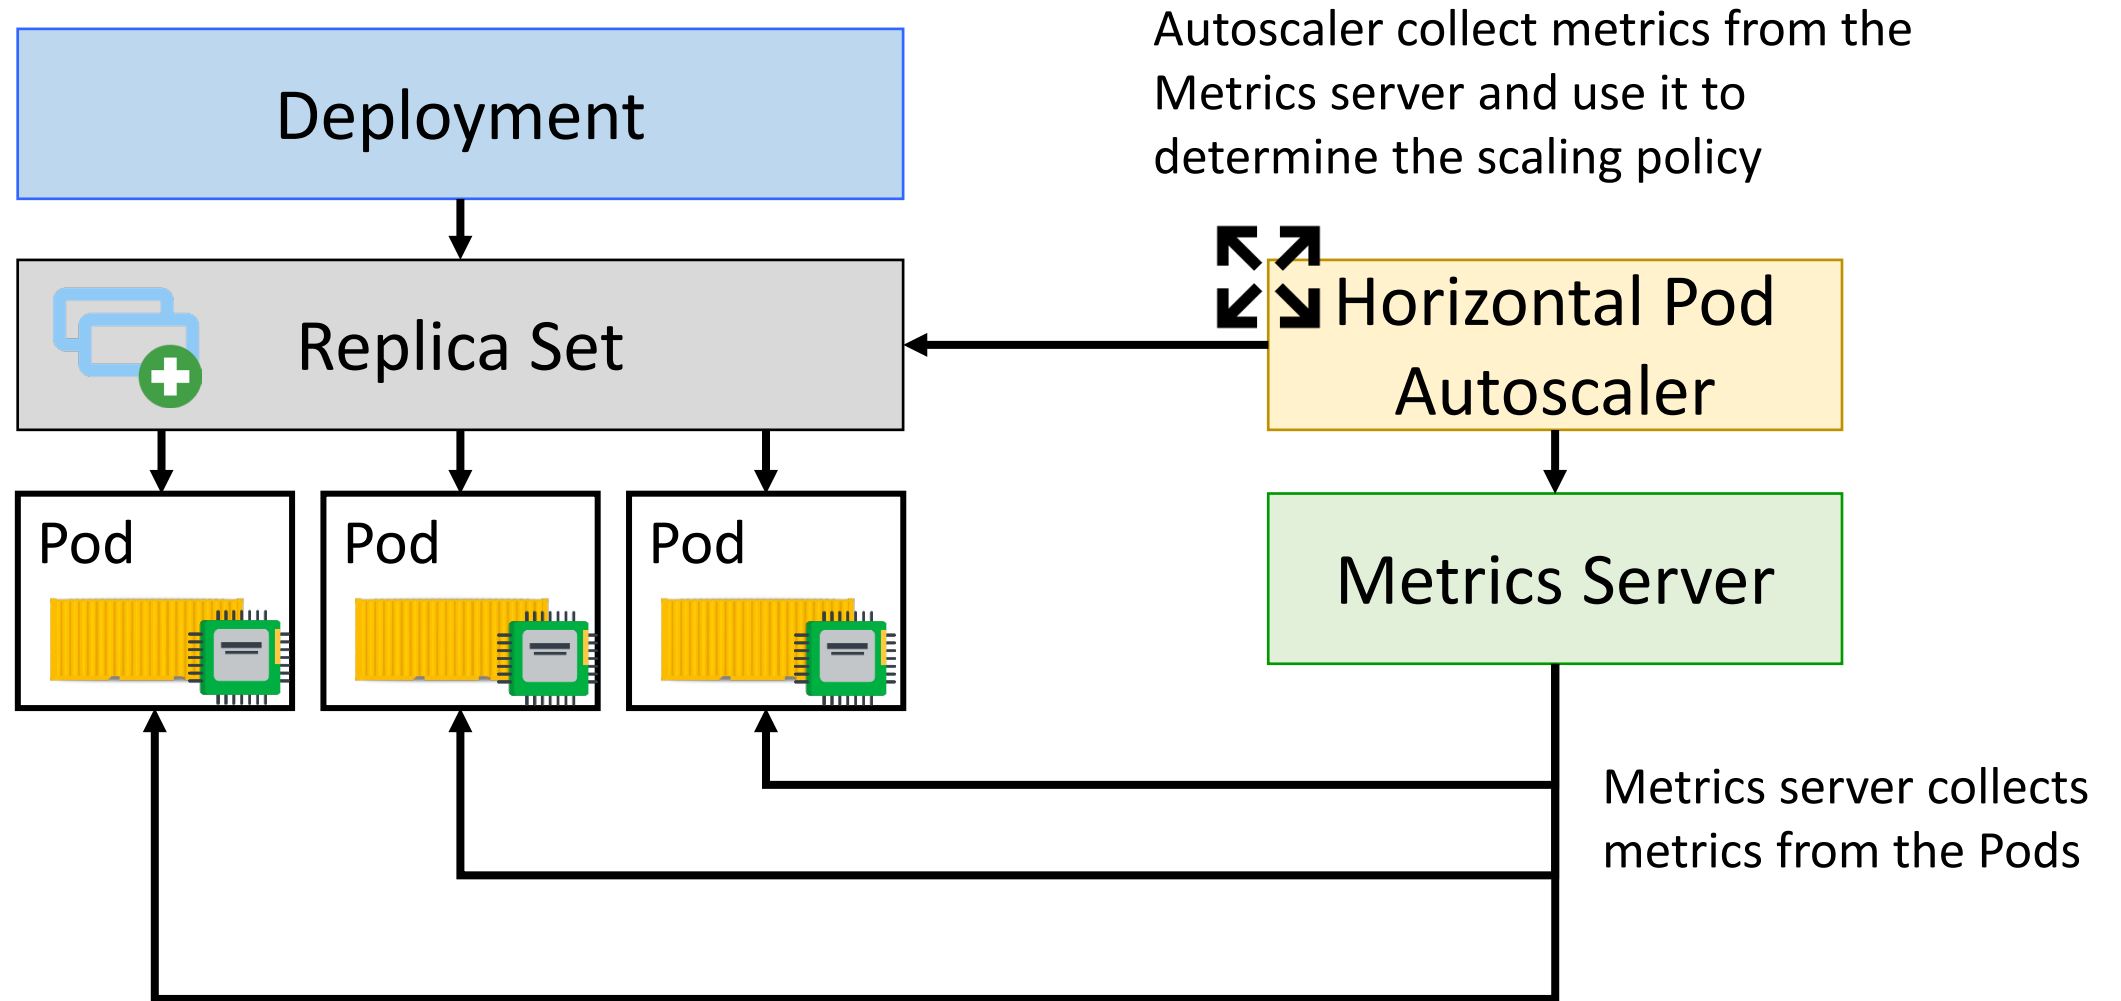


Setting Pod Request

- Horizontal scaler scales a Pod by determining if the Pod has breached a certain threshold
 - For memory and CPU
- Set the request for CPU and memory
 - Specify the minimum amount of compute resources required
- Resource type
 - CPU - measured in CPU units eg 100m is 100 millicores
 - 1 CPU in Kubernetes == 1 vCPU, Core, vCore, Hyperthread
 - Memory - 16M



Horizontal Pod Autoscaler





Requesting Resources

```
apiVersion: v1
kind: Pod
metadata:
  name: myapp
spec:
  containers:
  - name: myapp
    image: myapp:sha256:...
```

HPA only looks
at the CPU

```
resources:
```

```
requests:
```

```
  cpu: 100m
```

```
  memory: 16M
```

```
limits:
```

```
  memory: 32M
```

```
...
```

} Request the minimum amount of
compute resources

→ Describe the maximum amount
of compute resources required



Defining a Horizontal Pod Autoscaler

```
apiVersion: autoscaling/v2beta2
kind: HorizontalPodAutoScaler
metadata:
  name: myapp
spec:
  minReplicas: 1
  maxReplicas: 8
  scaleTargetRef:
    apiVersion: apps/v1
    kind: Deployment
    name: myapp
```

The deployment that
this HPA is targeting

Minimum and maximum number of replicas. Since the HPA is managing the replica set, this setting takes precedence over the deployment setting

metrics:

- type: Resource

resource:

name: cpu

target:

type: Utilization

averageUtilization: 80

- type: Resource

resource:



name: memory

target:

type: Utilization

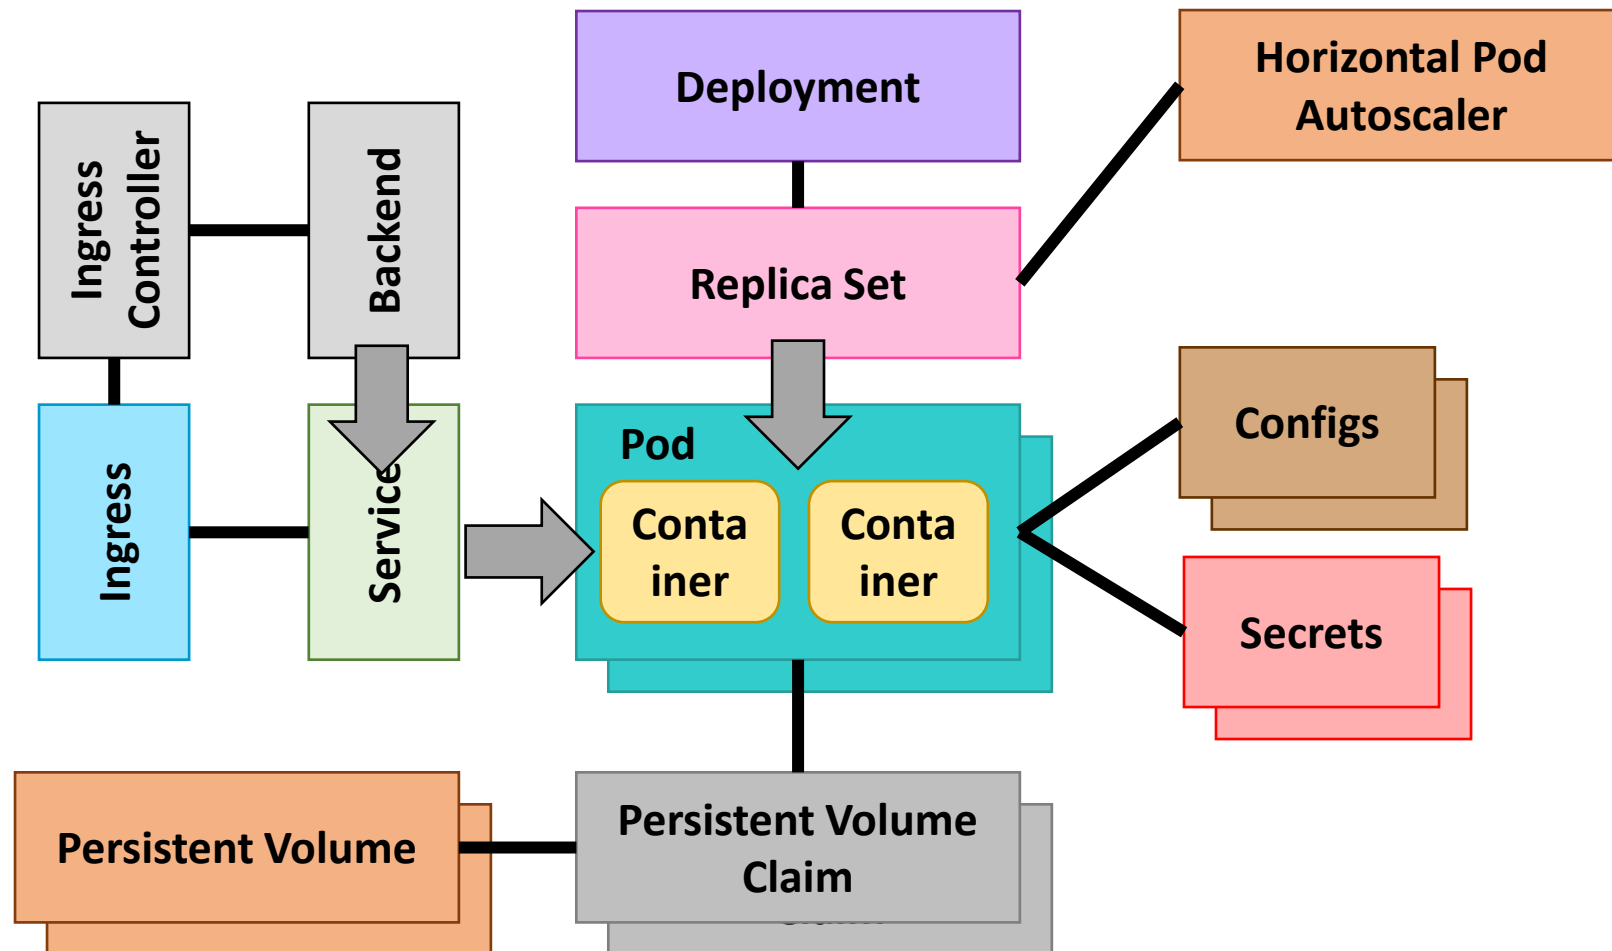
averageUtilization: 80

Scale the pod
when these
threshold are
breached





Kubernetes





Appendix



Using ConfigMaps

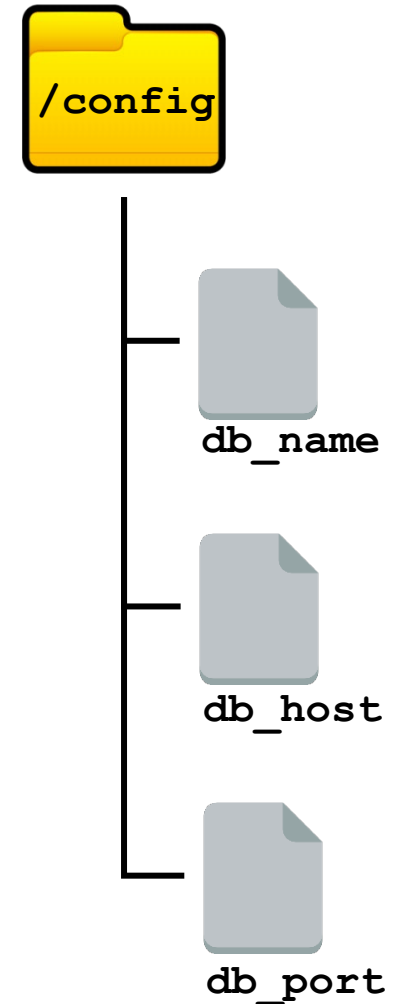
Injecting as environment variables

```
containers:
  ...
  env:
    - name: DB_NAME
      valueFrom:
        configMapKeyRef:
          name: myapp-config
          key: db_name
    - name: DB_HOST:
      valueFrom:
        configMapKeyRef:
          name: myapp-config
          key: db_host
```

Mounting as a volume

```
volumes:
  - name: config-volume
    configMap:
      name: myapp-config

containers:
  ...
  volumeMounts:
    - name: config-config
      mountPath: /config
```





Managing Context

- For grouping access parameters under a common name
 - Like a profile
 - Set the namespace, do not need the `-n` option

- Create a context

```
kubectl config set-context <context_name> --namespace=<name> \
    --cluster=<cluster_name> --user=<user_name>
```

- View current contexts

```
kubectl config view
```

- Use a context

```
kubectl config use-context <context_name>
```